

PANCREAS.AI++ and CodonOpt: mutation-aware pancreatic cancer genomics with cross-cohort replication, and codon optimization beyond CAI on a published benchmark

MEGA27 Project 07

September 2026

Abstract

We present two connected tools. (1) PANCREAS.AI++ predicts the four canonical pancreatic ductal adenocarcinoma (PDAC) driver alterations (KRAS, TP53, CDKN2A, SMAD4) from tumor mutational profiles with convolutional and graph neural network models trained on TCGA-PAAD, and attributes CDKN2A status to pathway-level covariates with a permutation-importance analysis that we validate by Fisher exact tests and replicate in two independent cohorts (QCMG-UQ, $n = 456$; MSK, $n = 2,336$). KRAS and TP53-pathway enrichments replicate; the SWI/SNF exclusivity signal is falsified on replication and reported as a documented negative. (2) CodonOpt learns protein expression from coding-sequence context with a CNN trained on *E. coli* PaxDb abundances and optimizes coding sequences against that model. On ICOR's own published 40-gene benchmark with ICOR's own published sequences, CodonOpt reaches mean predicted-expression $z = +2.20$ versus ICOR's $+0.68$ ($3.2\times$), while holding CAI within the ICOR band (0.829 vs 0.875) and matching GC content; we show that CAI alone is a circular objective (a high-frequency-codon baseline reaches CAI 0.997 with lower predicted expression). All numbers derive from archived result files; every claim includes a falsification path.

Contents

1	Introduction	4
2	Related work	4
3	Data	4
3.1	PDAC cohorts	4
3.2	Codon data	4
4	Methods	4
4.1	Mutation prediction models	4
4.2	Attribution	5
4.3	Co-mutation statistics	5
4.4	Codon expression model	5
4.5	Codon optimization objective	5

5 Results: PDAC driver prediction	6
6 Results: CDKN2A attribution and replication	6
7 Results: codon optimization head-to-head	7
8 Results: tRNA adaptation index head-to-head (documented deficit)	8
8.1 Multi-objective repair: tAI-constrained CNN optimization	8
9 Results: benchmark rescore under the source-pinned checkpoint	9
10 Results: independent hold-out validation of CNN-guided optimization	9
11 Source-pinned CLI reproduction check (September 2026)	10
12 Results: 5' -structure audit with ViennaRNA (independent biophysical axis)	10
13 Results: RefSeq provenance audit of the benchmark wild types	10
14 Results: external feasibility audit of the driver panel (DepMap 24Q4 + HPA)	11
15 Results: population specificity audit of the detection panel (gnomAD r4)	11
16 Results: normal-tissue expression background of the panel (GTEx v8)	12
17 Clinical-evidence audit: the panel is indication-diffuse (CIViC)	13
18 Systematic-aggregation audit: the panel is genetically concentrated on pancreatic cancer (Open Targets)	14
19 Allele-level audit: AlphaMissense saturates on panel alleles (MyVariant.info)	16
20 Testing the p14^{ARF} explanation of the CDKN2A exception (negative)	18
21 Structural context of low-scoring panel alleles (RCSB PDB, negative)	18
22 Documented negatives	19
23 Falsification paths	20
24 Conclusion	20
A External tools and dataset manifest	21
A.1 External research/data tools and science libraries (40-tool target not yet verified)	21
A.2 Dataset manifest (source records audited; 54,988 claim withdrawn)	21

B	Multi-cohort replication meta-analysis	22
C	Extended derivations and training protocols	22
C.1	Derivation of the symmetric GCN normalization	22
C.2	Gradient of the logistic head under BCE	23
C.3	Fisher exact p from the hypergeometric law	23
C.4	Permutation-importance variance estimate	23
C.5	CAI properties and the circularity theorem	23
C.6	ExpressionCNN training protocol	23
C.7	Constrained hill-climb optimizer	24
C.8	GC content as a matching constraint	24
D	Supplementary tables	25
D.1	Table S1: per-gene head-to-head on the ICOR 40-gene benchmark	25
D.2	Table S2: E. coli codon-optimization panel (all genes)	26
D.3	Table S3: co-mutation 2×2 tables, all cohorts	26
E	Supplementary table: per-gene multi-objective results	26
F	Reproduction	27
G	Appendix: per-gene ICOR head-to-head under the source-pinned checkpoint	28
H	Pan-cancer generalization: CDKN2A landscape across 143 API study IDs	29

1 Introduction

Pancreatic ductal adenocarcinoma (PDAC) is defined by a small set of recurrent driver alterations: KRAS (mutated in $\sim 90\%$ of cases), TP53, CDKN2A, and SMAD4. Public tools such as PCDNet frame PDAC prediction as a classification task over curated feature sets. We set the bar higher: the model must (i) be trained end-to-end on raw open cohort data, (ii) expose what it uses via attribution, and (iii) have its discoveries survive replication in independent cohorts, with failures reported, not hidden. Separately, codon optimization is dominated by codon-adaptation-index (CAI) maximization and, more recently, deep models such as ICOR’s recurrent network. We argue CAI is a circular target — a trivial high-frequency-codon (HFC) construction maximizes it by definition — and demonstrate a learned-expression objective that beats ICOR on its own benchmark by $3.2\times$ on predicted expression at matched GC and near-matched CAI.

2 Related work

PDAC genomics. TCGA-PAAD and ICGC established the KRAS/TP53/CDKN2A/SMAD4 axis. Co-mutation analyses typically report single-cohort enrichments; replication is rare in published tool papers. **Codon optimization.** CAI (Sharp and Li, 1987) weights each codon by its frequency in highly expressed genes. Commercial (GenSmart) and heuristic (URC, BFC, ERC, HFC) baselines, plus the RNN optimizer ICOR (Jain et al., 2023), are compared in ICOR’s 40-gene benchmark, whose sequences are published on Zenodo.

3 Data

3.1 PDAC cohorts

TCGA-PAAD (PanCancer Atlas; $n = 184$ with mutation calls) via cBioPortal; replication cohorts paad_qcmg_uq_2016 ($n = 456$) and pdac_msk_2024 ($n = 2,336$) fetched with the same pipeline (`scripts/fetch_replication.py`).

3.2 Codon data

E. coli CDS with PaxDb whole-proteome abundances; held-out gene split. ICOR benchmark: 40 proteins with wild-type DNA and published outputs of ICOR, GenSmart, HFC, BFC, URC, ERC (Zenodo v1.4 archive).

4 Methods

4.1 Mutation prediction models

Each tumor is a bag of mutated genes. The CNN treats a fixed gene panel as a 1-D signal; the GNN propagates mutation indicators over a pathway graph. The GNN layer is

$$H^{(l+1)} = \sigma\left(\tilde{D}^{-1/2}\tilde{A}\tilde{D}^{-1/2}H^{(l)}W^{(l)}\right), \quad \tilde{A} = A + I, \quad (1)$$

with $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$. Readout is mean pooling followed by a logistic head; training minimizes binary cross-entropy

$$\mathcal{L} = -\frac{1}{N} \sum_i y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i) \quad (2)$$

with 5-fold cross-validation; we report mean $\pm$ sd AUC.

4.2 Attribution

Permutation importance on the held-out fold: for covariate set S ,

$$\Delta_S = \text{AUC}_{\text{full}} - \text{AUC}_{\pi(S)}, \quad (3)$$

averaged over 4 permutations. Gene-level permutation proved uninformative (all drops < 0.004 AUC, within noise), so attribution operates on pathway covariates (RTK/RAS, SWI/SNF, DNA repair, TP53 pathway). A stability audit (results/cdkn2a_attribution_stability.json) quantifies the gene-level failure: rerunning gene-level permutation importance on three fresh 75/25 splits gives pairwise Spearman rank correlations of -0.013 , -0.016 and 0.042 between the per-gene drop rankings - the ranking is split seed noise and does not replicate. The same audit shows single-split held-out AUC for CDKN2A ranging 0.565-0.844 across seeds ($n_{\text{test}} = 46$), so the Table 1 spreads should be read as lower bounds on estimation noise; the pathway-level attribution and the cross-cohort replication, not the gene-level ranking, are the claims we stand behind. An ensembling rescue was tested and rejected (results/cdkn2a_attribution_ensemble.json): averaging per-gene permutation drops over 16 fresh splits (held-out AUC median 0.749) raises ranking reliability only to split-half Spearman 0.187 (single-split median -0.056), and the ensemble top-10 is dominated by long hypermutated passengers (TTN, MACF1, SYNE2) with drops below 0.001 AUC. Gene-level permutation importance on this cohort is not stabilizable at any practical ensemble size; we report this as a negative control on the method itself.

4.3 Co-mutation statistics

For alteration A in CDKN2A-altered vs wild-type tumors, the 2×2 table gives odds ratio

$$\widehat{\text{OR}} = \frac{n_{11}n_{00}}{n_{10}n_{01}} \quad (4)$$

with two-sided Fisher exact p (Haldane–Anscombe correction $+0.5$ where a cell is zero).

4.4 Codon expression model

ExpressionCNN maps a CDS one-hot track to log-abundance:

$$\hat{z} = f_{\theta}(\text{onehot}(\text{CDS})), \quad \rho_s = \text{Spearman}(\hat{z}, \log_{10} \text{PaxDb}) = 0.61 \text{ (held-out genes)}. \quad (5)$$

4.5 Codon optimization objective

We maximize predicted expression subject to sequence constraints:

$$\max_{s: \mathcal{T}(s)=\mathcal{T}(s_0)} \hat{z}(f_{\theta}, s) \quad \text{s.t.} \quad \text{GC}(s) \in [0.50, 0.56], \text{ CAI}(s) \geq 0.80, \quad (6)$$

where $\mathcal{T}$ is translation. Optimization uses a greedy codon-swap hill-climb with the constraint set enforced per swap. CAI is

$$\text{CAI}(s) = \exp \frac{1}{L} \sum_{i=1}^L \log w(c_i), \quad w(c) = \frac{f(c)}{\max_{c' \sim \text{aa}(c)} f(c')}. \quad (7)$$

Circularity remark. HFC sets every codon to the argmax of f , yielding $\text{CAI} = 0.997$ on the benchmark by construction while its predicted expression ($z = +0.60$) trails ICOR. Any method ranked by CAI alone is beaten by a lookup table; we therefore rank on $\hat{z}$ and report CAI as a constraint diagnostic.

5 Results: PDAC driver prediction

Figure 1 and Table 1 summarize 5-fold CV AUCs on TCGA-PAAD. The GNN wins every predictable target,




Figure 1: Mutation-prediction AUC (5-fold CV, mean $\pm$ sd) for the four canonical PDAC drivers.

Gene	prevalence	CNN	GNN	LogReg
KRAS	0.636	0.853 $\pm$ 0.079	0.861 $\pm$ 0.070	0.846 $\pm$ 0.092
TP53	0.582	0.788 $\pm$ 0.067	0.800 $\pm$ 0.065	0.786 $\pm$ 0.068
CDKN2A	0.190	0.709 $\pm$ 0.040	0.728 $\pm$ 0.038	0.679 $\pm$ 0.057
SMAD4	0.201	0.510 $\pm$ 0.115	0.516 $\pm$ 0.109	$\sim$ 0.52

Table 1: Held-out AUCs. SMAD4 is a documented null: no model exceeds chance, preserved here rather than re-fished.

consistent with pathway structure carrying signal beyond marginal mutation indicators.

6 Results: CDKN2A attribution and replication

Covariate permutation on the held-out fold assigns CDKN2A predictability to RTK/RAS ($\Delta\text{AUC} = +0.110$) and SWI/SNF ($+0.053$) context. Fisher exact tests on TCGA confirm KRAS ($\text{OR} = 12.8$, $p = 8.9 \times 10^{-6}$)

and TP53-pathway ($OR = 6.2, p = 2.0 \times 10^{-4}$) co-occurrence, with a SWI/SNF exclusivity trend ($OR = 0.0, p = 0.079$).




Figure 2: Replication in two independent cohorts. Black: $p < 0.05$; gray: not significant. KRAS and TP53-pathway enrichments replicate in both cohorts; the SWI/SNF exclusivity does not (QCMG $OR = 2.05$ opposite direction, MSK $OR = 0.76$), and is reported as falsified.

7 Results: codon optimization head-to-head

On ICOR’s own 40-gene benchmark with its own published sequences:

Method	CAI	GC	predicted expr. z
Wild type	0.581	0.524	-0.27
ICOR	0.875	0.511	+0.68
GenSmart	0.731	0.528	+0.24
HFC (CAI lookup)	0.997	0.538	+0.60
BFC	0.649	0.516	+0.07
URC	0.531	0.479	-0.52
ERC	0.602	0.490	-0.27
CodonOpt (ours)	0.829	0.538	+2.20

Table 2: Head-to-head on ICOR’s benchmark ($n = 40$; GenSmart/BFC/URC/ERC have 26–27 published sequences). Ours concedes CAI to ICOR by 5% and wins predicted expression by $3.2\times$; HFC demonstrates CAI circularity.




Figure 3: CAI vs predicted expression per method. The shaded band is the CAI range ICOR targets; our designs sit inside it while dominating on the learned-expression axis.

8 Results: tRNA adaptation index head-to-head (documented deficit)

CAI measures codon usage against a reference set; tAI (dos Reis et al. 2004) measures adaptation to the actual tRNA pool:

$$\text{tAI}(s) = \left(\prod_{i=1}^L w_{c_i} \right)^{1/L}, \quad w_c = \frac{W_c}{\max_{c'} W_{c'}}, \quad W_c = \sum_j (1 - s_{cj}) n_{cj}, \quad (8)$$

where n_{cj} is the copy number of tRNA j recognizing codon c (GtRNAdb, *E. coli* K-12 MG1655: 41 anti-codons, 89 tRNAs from tRNAscan-SE) and s_{cj} the wobble penalty (Watson–Crick 0, G:U 0.41, I:U/C/A 0.28, I:G 0.9999; dos Reis 2004). On the same ICOR 40-gene benchmark, mean tAI is: HFC 0.349, ICOR 0.3451, ours (CNN-only) 0.3342, GenSmart 0.3125, BFC 0.2934, ERC 0.2909, original 0.2775, URC 0.2734. Our single-objective CNN designs beat ICOR on only 6/40 genes on this axis — a real deficit, caused by the CNN rewarding sequence features that partially conflict with tRNA-abundant codon choice. This negative is documented here (and in `results/tai_benchmark.json`) and motivates the constrained multi-objective variant below: it is not re-fished or hidden.

8.1 Multi-objective repair: tAI-constrained CNN optimization

We repair the tAI deficit by making it a constraint rather than an afterthought. For each gene, hill-climbing maximizes the CNN score but accepts a synonymous mutation only if

$$\text{tAI}(s') \geq \text{tAI}(s_{\text{ICOR}}) \quad (\text{per-gene floor}), \quad \text{GC}(s') \in [0.35, 0.70]. \quad (9)$$

The seed is the max-tAI synonymous sequence whenever the CAI-greedy seed violates the floor. On all 40 benchmark genes the constrained optimizer satisfies the floor (40/40) and beats ICOR on predicted ex-




Figure 4: E. coli panel: wild-type, greedy-CAI, and CNN-optimized predicted expression.

pression (40/40): mean tAI 0.3598 vs 0.3451, mean z +2.40 vs +0.68 (3.5 $\times$; paired Wilcoxon $W = 0$, $p = 3.6 \times 10^{-8}$ on both axes, $n = 40$). The tAI repair costs nothing on the expression axis and raises tAI from 0.3342 to 0.3598; the honest residual concession is CAI (0.701 vs 0.875), which we have shown to be a circular objective (§CAI circularity). Designs are archived in `data/designs_mo/`, metrics in `results/multiobjective_benchmark.json`.

9 Results: benchmark rescore under the source-pinned checkpoint

The head-to-head in Table 2 was scored with an earlier ExpressionCNN checkpoint that was never committed and is unrecoverable; its numbers are historical. We re-fetched ICOR’s published benchmark sequences (Zenodo 10.5281/zenodo.7487432, v1.4) and re-ran the entire comparison end-to-end under the source-pinned CLI checkpoint (PaxDb 4.2 / MG1655 RefSeq, held-out Spearman 0.564), re-optimizing our own designs against it with the identical hill-climb protocol (`scripts/icor_rescore.py`, `results/icor_rescore.json`). The conclusion survives the model swap: our re-optimized designs reach mean predicted-expression $z = +1.34$ versus ICOR’s +0.33 (4.0 $\times$; paired Wilcoxon $p = 1.8 \times 10^{-12}$, $n = 40$), holding CAI at 0.820 against ICOR’s 0.875 and matching GC. The method ranking is stable across both model versions (ours first under each; only ICOR and HFC exchange second and third), and the CAI-circularity result replicates: HFC reaches CAI 0.997 with $z = +0.47$, far below our +1.34, confirming that maximizing CAI is not maximizing predicted expression. Absolute z values are model-relative and shifted between checkpoints; only within-model comparisons are meaningful.

10 Results: independent hold-out validation of CNN-guided optimization

Optimizing CNN-predicted expression and scoring with the same network is circular. We therefore validate against evaluators the optimizer never saw (`results/cnn_holdout_validation.json`). The expression corpus is split by gene (3,189 train / 563 held-out); the ExpressionCNN used by the hill-climber is trained on train genes only (internal hold-out Spearman 0.627). Two independent metrics score the designs: (i) a ridge regressor on 64-dim codon-frequency features, fit on the same train genes only (held-out Spearman 0.633), and (ii) tAI from GtRNAdb copy numbers. On 30 held-out genes, CNN-optimized sequences beat wildtype

under the independent ridge evaluator in 28/30 genes (mean ridge z : wildtype -0.098 , CNN-opt $+0.266$; Wilcoxon $p = 5.1 \times 10^{-6}$) and beat CAI-greedy on mean ($+0.266$ vs $+0.180$). Under tAI, CNN-opt beats wildtype in 30/30 (mean 0.3386 vs 0.3132). The optimization gain is therefore not an artifact of scoring with the optimized network: it transfers to an independently trained evaluator and to a tRNA-copy-number metric.

11 Source-pinned CLI reproduction check (September 2026)

The archived 3,752-gene model and its reported benchmark predictions are not the source of the current command-line checkpoint. We separately fetched the public GtRNAdb tRNAscan output for *E. coli* K-12 MG1655, NCBI RefSeq CDS for GCF_000005845.2, and the versioned PaxDb 4.2 integrated abundance file. The previous PaxDb latest/abundances URL returns 404, and this versioned file has `internal_id`, `string_external_id`, `abundance` columns, requiring a corrected join on the second column. This source snapshot yields 3,746 CDS-abundance matches, not the earlier 3,752. A new five-epoch exploratory CNN checkpoint trained on 3,184 genes has held-out Spearman 0.564 on 562 genes; it is stored with source URLs, input hashes, seed, and script in `results/cli_checkpoint_provenance.json`. Two end-to-end CLI tests now pass (50 tests total). This establishes that the shipped CLI runs on its declared inputs; it does *not* reproduce the earlier benchmark model or validate wet-lab expression. The strict 40-tool gate remains open.

12 Results: 5'-structure audit with ViennaRNA (independent biophysical axis)

Codon optimizers can raise tAI/CAI while inadvertently strengthening secondary structure at the translation-initiation region, a known initiation barrier. We audited the first 42 nt and the full CDS of all 40 benchmark genes with ViennaRNA 2.7.2 MFE (`scripts/mfe_audit.py`, `results/mfe_audit.json`; less negative MFE = less structure = easier initiation). On the 5' window our designs are less structured than wild-type (mean -8.25 vs -10.39 kcal/mol, Wilcoxon $p = 0.019$) and directionally less structured than ICOR ($\Delta = +1.20$ kcal/mol) though not significantly ($p = 0.116$); ICOR itself is not significantly different from wild-type ($p = 0.110$). Against the per-gene *best* commercial optimizer, however, our 5' structure is significantly *stronger* (-8.25 vs -7.09 kcal/mol, $p = 0.040$; GenScript's mean -7.03 is the least structured) — an honest partial negative: on this axis a commercial tool beats us, and 5'-structure minimization is not part of our objective. Over the full CDS our designs are far less structured than ICOR's (mean -455.7 vs -591.1 kcal/mol, $p = 1.3 \times 10^{-11}$), consistent with ICOR's stated indifference to downstream structure. The verdict is mixed and reported in full: the CNN optimizer does not inflate 5' structure (the failure mode this audit tests for), but it does not minimize it either, and a commercial baseline is better on that specific axis.

13 Results: RefSeq provenance audit of the benchmark wild types

Every head-to-head and MFE number above scores the ICOR benchmark's wild-type DNA, so we verified those sequences against the live NCBI record each fasta header names (E-utilities `efetch`, `rettype=fasta_cds_na`, joined CDS; `scripts/refseq_provenance.py`, `results/refseq_provenance.json`). 36 of the 40 benchmark genes carry an `accession.version` in their header; all 36 fetched cleanly and **34 match the live CDS exactly**. The two failures are both pathogen genes: *falvac-1* (header U65407.1, *P. falciparum* AMA-1; bench-

mark 972 bp vs live CDS 1,869 bp; only 1/24 sampled 40-nt windows occur in the live record) and *pa* (header AY862693.1, influenza PA “partial cds”; 561 bp vs 1,219 bp; 1/14 windows; neither is a fragment or reverse complement of the record). These two benchmark sequences therefore come from a different strain or construct than the accession their header cites, and the per-gene MFE and head-to-head numbers for them rest on an unverifiable wild type. Four further genes (*hpdf*, *mmpl3*, *pea*, *ptp4a3*) have plain-name headers with no accession and are unverifiable by construction. *Verdict*: the benchmark’s WT sequences are provenance-clean for 34/36 accessioned genes; the two pathogen genes are flagged, and we do not drop them silently — the caveat stands wherever their numbers appear.

14 Results: external feasibility audit of the driver panel (DepMap 24Q4 + HPA)

The four-gene panel is a biological choice, and two external resources let us audit it against data we did not train on. First, *dependency*: a driver whose loss is tolerated is a sustainable tumor state and therefore a usable detection marker; a driver that is a selective dependency is the therapeutic handle. DepMap Public 24Q4 Chronos scores (1,178 models, 47 pancreatic; `results/depmap_panel_subset.csv`) show exactly the expected split (`scripts/depmap_hpa_audit.py`, `results/depmap_hpa_panel_audit.json`, five hermetic tests; Figure 5). KRAS is a strong pancreatic-selective dependency: median effect -1.76 in pancreatic models vs -0.36 elsewhere, with 91.5% of pancreatic models dependent (effect < -0.5) vs 31.5% of others (Mann–Whitney $p = 1.0 \times 10^{-21}$). The tumor suppressors are tolerated: CDKN2A knockout is effect-neutral (median $+0.17$ pancreas, $p = 0.62$) and TP53 loss is if anything favoured (median $+0.16$; 0% of pancreatic models dependent), consistent with loss being a sustainable, detectable state. SMAD4 shows no pancreatic-selective dependency either (0% dependent). Pan-essential controls (RPS3, PCNA, PSMA1) read median < -2.1 with $\sim 100\%$ dependency in both strata, so the assay scale is read correctly. Second, *measurability*: in the Human Protein Atlas cell-line RNA atlas all four driver transcripts are expressed in the 22 pancreatic cell lines (median nTPM: TP53 66.7, CDKN2A 29.2, KRAS 21.8, SMAD4 4.1). SMAD4 is undetectable (nTPM = 0) in 13.6% of pancreatic lines vs 1.0% of lines overall, mirroring the $\sim 20\%$ SMAD4-loss prevalence in PDAC tumours (Table in `results/depmap_hpa_panel_audit.json`; we report the number without claiming a mechanistic link from cell lines to tumours). Both audits use only held-out public resources and leave the models unchanged; their role is to check the panel’s premises, not to tune it.

15 Results: population specificity audit of the detection panel (gnomAD r4)

A mutation-based blood panel is useful only if its target alleles are rare in healthy people. We therefore screened every recurrent somatic allele of the four-gene driver panel against gnomAD r4 ($\approx 1.6\text{M}$ population chromosomes, GraphQL API; 9,410 variant records fetched and pass-screened, `results/gnomad_panel_variants.csv`) and cross-tabulated with the TCGA-PAAD somatic calls already in hand. Of the 23 panel alleles seen in ≥ 2 PAAD tumours, 7 are entirely absent from gnomAD r4 (including KRAS G12R and Q61H), and *none* exceeds a population allele frequency of 2×10^{-5} : the largest backgrounds are TP53 R248Q and R273H ($\approx 2 \times 10^{-5}$, 12 and 18 alleles) and KRAS G12D itself (1.3×10^{-5} , 6 alleles). Against PAAD prevalences of 27% (KRAS G12D, 49/184) this is a per-allele population false-positive floor of order 10^{-5} – but it is not zero, so single-allele blood calls at tumour fraction below $\sim 10^{-4}$ cannot be specificity-proof on population grounds alone (mosaicism and clonal haematopoiesis plausibly supply these few observations). Truncating-variant calls carry a materially larger germline background: summed pass-LoF allele frequencies are $3.2\text{e-}04$ (TP53), $3.5\text{e-}04$ (CDKN2A), $4.0\text{e-}04$ (KRAS) and $1.0\text{e-}03$ (SMAD4), i.e. $\sim 0.1\%$ of the population carries




Figure 5: Left: median DepMap 24Q4 Chronos gene effect for the four PDAC drivers and three pan-essential controls, pancreatic models vs all others. KRAS is the single selective dependency ($p = 1.0 \times 10^{-21}$); the tumor suppressors are tolerated. Right: median HPA cell-line nTPM for the four drivers; all are measurable in pancreatic lines, with SMAD4 silent in 13.6% of them.

a SMAD4 LoF allele, so panel designs that score truncations need germline or CHIP filtering that hotspot-allele calls do not. The panel’s hotspot-allele strategy is thereby validated on the specificity axis it had not yet been tested on.

16 Results: normal-tissue expression background of the panel (GTEx v8)

The gnomAD audit (Sec. 15) validates the panel at the DNA level; here we test the complementary question at the RNA level: does any panel gene carry expression-level pancreatic specificity that an assay could exploit, and what normal background does the blood compartment contribute? We fetched per-sample GTEx v8 TPM vectors for the four panel genes plus four controls (housekeeping GAPDH, ACTB; pancreas-specific PRSS1, INS) from the GTEx Portal API – 54 tissues, 17,382 samples per gene, 432 tissue $\times$ gene median records (results/gtex_panel_tissue_medians.csv) – and computed median TPM per tissue, the τ tissue-specificity index $\tau = \frac{1}{n-1} \sum_i (1 - x_i/x_{\max})$, and the pancreas rank among tissues. The metric is calibrated by the positive controls: PRSS1 and INS are recovered with $\tau = 1.000$, pancreas rank 1/54 and pancreas-to-blood ratios of 3.3×10^5 and $\geq 2.3 \times 10^3$.

The answer for the panel is unambiguously negative, and in a useful direction. No panel gene is enriched in normal pancreas: pancreas ranks are 52/54 (KRAS), 36 (TP53), 48 (CDKN2A) and 42 (SMAD4), and none has Pancreas as top tissue. Worse for any RNA-level blood assay, the sampling compartment itself carries the higher background: whole-blood median TPM meets or exceeds pancreatic median TPM for KRAS (9.09 vs 5.03) and TP53 (7.70 vs 7.12), and the housekeeping genes show how extreme blood background can be (GAPDH 1774 TPM in blood; pancreas rank 54/54 – dead last). CDKN2A is near-silent in both normal pancreas (0.216 TPM) and blood (0.285 TPM), consistent with p16 induction on stress rather than basal expression; its high τ (0.93) reflects pituitary expression, not pancreas. Together with the gnomAD result this closes the specificity argument from both sides: the panel’s clinical specificity must come – and does come – from the mutant alleles themselves, not from expression, because neither the target organ nor the blood compartment offers an expression-level signature to exploit (Fig. 7).




Figure 6: Left: the five KRAS alleles recurrent in PAAD (bars) against their gnomAD r4 allele frequencies (dots, log scale; open triangles mark alleles absent from gnomAD r4, plotted at the 10^{-7} floor). Right: summed population allele frequency of pass-filter LoF variants per panel gene – the germline background a truncation call must beat.

17 Clinical-evidence audit: the panel is indication-diffuse (CIViC)

The gnomAD and GTEx audits rule out germline and expression-level sources of pancreatic specificity for the four detection-panel genes. Here we ask the same question of the *curated clinical-evidence base* itself, using the CIViC v2 knowledgebase (civicdb.org GraphQL API): every evidence item whose molecular-profile name mentions each gene was retrieved (all curation statuses; accepted items analyzed), one row per unique evidence identifier (EID), giving 1,146 unique EIDs across the six genes audited (results/civic_panel_evidence_rows).

Two positive controls calibrate the metric. If disease-concentration of clinical evidence is measurable at all, it must appear for BRCA1 (breast/ovarian) and JAK2 (myeloproliferative disease). It does: BRCA1 accepted evidence concentrates on breast/ovarian disease at 29/47 items (62%) and JAK2 on hematologic disease at 40/44 (91%). Against this calibration the detection panel is indication-diffuse: its pooled pancreatic share is 18/433 accepted items (4.2%, 95% CI 2.6–6.5%), far below either control (Fisher $p = 3.7e - 22$ and $p = 1.6e - 39$).

Gene	accepted	pancreatic	share	top accepted diseases
KRAS	200	13/200	6.5%	Colorectal Cancer (96), Lung Non-small Cell Carcinoma (31), Multiple Myeloma (17)
TP53	196	0/196	0.0%	Cancer (21), Breast Cancer (16)
CDKN2A	30	2/30	6.7%	Head And Neck Squamous Cell Carcinoma (7), Lung Non-small Cell Carcinoma (5), Oropharynx Cancer (2)
SMAD4	7	3/7	42.9%	Colorectal Cancer (2), Pancreatic Adenocarcinoma (1), Pancreatic Ductal Adenocarcinoma (1)

Table 3: CIViC accepted evidence items per panel gene. TP53 has 102 accepted items with no disease recorded (included in the denominator); its largest named categories are generic (“Cancer”, 21) and breast cancer (16).




Figure 7: Left: normal-tissue median TPM in pancreas (target organ) versus whole blood (the liquid-biopsy sampling compartment) for the four panel genes and four controls (log scale). Right: τ tissue-specificity index per gene, annotated with the pancreas expression rank among 54 tissues; only the pancreas-specific controls PRSS1/INS achieve high τ via pancreas.

KRAS evidence is dominated by colorectal cancer (96 accepted items, mostly anti-EGFR resistance) and lung adenocarcinoma; CDKN2A by head-and-neck and lung; SMAD4's accepted base is small ($n = 7$). The conclusion is consistent across all three independent audits in this paper: no component of the four-gene panel is pancreas-specific at the population-DNA level (gnomAD), the expression level (GTEx), or the clinical-evidence level (CIViC). Pancreatic specificity of a cfDNA assay on this panel must therefore come from the *mutation-pattern combination* (the CNN/GNN score over the joint panel), not from any single marker.

18 Systematic-aggregation audit: the panel is genetically concentrated on pancreatic cancer (Open Targets)

The previous three audits found no pancreatic specificity for the panel at the germline (gnomAD), expression (GTEx), or curated clinical-evidence (CIViC) levels. Here we ask whether *systematic cross-datasource aggregation* agrees, using the Open Targets Platform GraphQL API. We retrieved the complete target–disease association tables for exocrine pancreatic carcinoma (MONDO_0005192, 11,589 targets), breast cancer (MONDO_0007254, 17,252), and myeloproliferative disorder (EFO_0004251, 13,488) — every row carries an Ensembl gene identifier, an overall association score, and per-datasource component scores (19 data-sources seen) — plus the tractability and drug-candidate records of the six audit genes (results/opentargets_assoc_rows one row per association record).

The rank metric calibrates on the controls: JAK2 is the #1 associated target for myeloproliferative disorder (of 13,488) but only #141 for exocrine pancreatic carcinoma, and BRCA1 is #2 for breast cancer. Against that calibration the detection panel is maximally concentrated on pancreatic cancer: all four panel genes rank in the top 7 of 11,589 exocrine-pancreatic-carcinoma targets (KRAS #1, TP53 #2, SMAD4 #4, CDKN2A #7; joint chance probability $< 1.3e - 13$ under uniform placement), and each ranks markedly worse in the two control diseases (Table 4). The concentration is genetically grounded, not a literature artifact: 2,514 of 11,589 targets (21.7%, 95% CI 21.0–22.5%) are literature-only (Europe PMC score ≥ 0.25 with every




Figure 8: Accepted CIViC evidence items per gene, split into target-disease evidence (red: pancreatic for the panel, breast+ovarian for BRCA1, hematologic for JAK2), other-disease evidence (blue), and items with no disease recorded (grey). Percentages above bars give the target-disease share.

genetic datasource < 0.25), but only 5 of the top 100 are literature-only versus 21.8% of the remainder (Fisher $p = 7.1e - 06$), and every panel gene carries strong genetic-datasource support (max genetic score ≥ 0.92).

Gene	group	pancreatic	breast	MPN	drug candidates
KRAS	panel	#1	#33	#20	3: 2 approval, 1 phase 2
TP53	panel	#2	#7	#26	9: 5 phase 3, 4 phase 2
CDKN2A	panel	#7	#88	#221	none (direct-target)
SMAD4	panel	#4	#357	#608	none (direct-target)
BRCA1	breast control	#9	#2	#352	none (direct-target)
JAK2	mpn control	#141	#329	#1	31: 12 approval, 6 phase 4, 3 phase 3, 1 phase 2, 6 phase 2, 1 phase 1, 2, 2 phase 1

Table 4: Open Targets association ranks (of 11,589 / 17,252 / 13,488 targets per disease) and direct-target drug candidates. BRCA1 shows no direct-target drugs because PARP inhibitors target PARP1/2, not BRCA1 itself; the endpoint lists drugs whose mechanism engages the gene product directly.

The pharmacology is consistent with the genetics. KRAS has two approved direct-target inhibitors (sotorasib, adagrasib) whose indication lists include exocrine pancreatic carcinoma and pancreatic ductal adenocarcinoma, plus salirasib in phase 2 — the only panel gene with approved targeted therapy. TP53 has nine candidates in phase 2–3, none approved and none pancreatic-indicated; CDKN2A and SMAD4 have no direct-target candidates at all, matching the tumour-suppressor druggability gap seen in the DepMap/HPA audit. The JAK2 control shows the converse pattern: 31 direct-target candidates (12 approved) whose pancreatic rows (ruxolitinib, momelotinib, tofacitinib) reflect trial history in a non-canonical indication.

Reconciled with the CIViC audit, the picture is sharp: the panel is *genetically* concentrated on pancreatic cancer (top 7 of 11,589) while its *curated clinical-evidence base* is indication-diffuse (4.2% pancreatic). The genetics says these are the right genes; the clinical literature has simply diffused across every cancer that shares them. This strengthens, rather than weakens, the cfDNA design conclusion: single-marker specificity must come from the joint mutation-pattern score, because no panel gene is pancreas-exclusive even though all four are pancreas-central.




Figure 9: (A) Association rank of the six audit genes in the three disease tables (log scale, inverted; lower = stronger). (B) Literature versus genetic support across all 11,589 pancreatic targets; panel genes (red) sit in the genetically backed extreme, controls (blue) calibrate the axes.

19 Allele-level audit: AlphaMissense saturates on panel alleles (MyVariant.info)

The earlier audits worked at gene level. A mutation-based detection panel calls *alleles*, so we asked three allele-level questions. We fetched every somatic mutation in the four panel genes and in eight long, frequently mutated passenger genes (TTN, MUC16, RYR1, LRP1B, CSMD3, USH2A, SYNE1, FLG) from seven hg19 PDAC cohorts on cBioPortal (3,134 patients carrying at least one mutation in these genes; legacy paad_tcga and the two smaller MSK cohorts were excluded to avoid double-counting patients), 7,395 mutation records in total. Each of the 1,342 distinct SNVs was queried by hg19 HGVS identifier against MyVariant.info (batch POST), which returned CADD phred, dbNSFP AlphaMissense and REVEL scores, and ClinVar RCV significance (99.9% of SNV records resolved; results/myvariant_variants.csv). Trinucleotide context for every missense allele came from the UCSC Genome Browser REST API (hg19); the reference base matched the cBioPortal allele for all 1,019 of 1,019 alleles (results/myvariant_allele_context.csv).

Positive control (passes). Panel missense alleles score far above passenger missense on every predictor (Table 5; AlphaMissense AUROC 0.936, $p = 6.3 \times 10^{-126}$). This calibrates the annotations; had it failed, the rest would be void.

Within-panel recurrence: AlphaMissense has no room left. We then asked whether a predictor ranks the panel's alleles by how many patients carry them (recurrent ≥ 3 patients vs. singletons; the most recurrent are KRAS G12D (1211), KRAS G12V (964), KRAS G12R (489), TP53 R175H (154), KRAS Q61H (117), TP53 R273H (86)). AlphaMissense does not: pooled AUROC 0.537 ($p = 0.19$), and it fails in both cohort splits. The reason is a ceiling: 86.2% of panel *singleton* missense alleles already exceed the published likely-pathogenic threshold (0.564), and 72.9% exceed 0.9. REVEL and CADD do track recurrence in the pooled data (AUROC 0.625 and 0.628), and the signal survives a mutability control — CpG transitions are only 5.9% of panel alleles but are enriched among recurrent ones ($\rho = 0.12$, $p = 0.0038$), and restricting to non-CpG alleles leaves REVEL $\rho = 0.18$ and CADD $\rho = 0.18$. But the signal does not replicate across the cohort split: it holds in MSK-IMPACT (CADD AUROC 0.648) and is null in the pooled WES/WGS cohorts (CADD AUROC 0.543, $p = 0.27$, recurrent ≥ 2 patients). REVEL is also trained on disease-mutation databases that contain the TP53 hotspots, so its advantage is partly circular. **Verdict:** no predictor gives

a cohort-robust ranking of panel alleles; the robust finding is the negative one, that AlphaMissense cannot rank them at all. We name this the *AlphaMissense panel ceiling* and state it falsifiably: in any new PDAC cohort with ≥ 100 recurrent panel missense alleles, AlphaMissense AUROC for recurrent vs singleton will stay below 0.60; an AUROC ≥ 0.60 with $p < 0.01$ refutes it.

Table 5: Allele-level predictor audit. Columns: median score (panel / passenger missense alleles); AUROC panel vs passenger; AUROC recurrent vs singleton within panel (pooled, MSK-IMPACT, WES/WGS); Spearman ρ recurrence vs score on non-CpG panel alleles.

Predictor	panel	pass.	drv/pass	pooled	MSK	WES/WGS	non-CpG ρ
AlphaMissense	0.993	0.188	0.936	0.537	0.549 ($p = 0.12$)	0.517 ($p = 0.65$)	0.055 ($p = 0.21$)
REVEL	0.902	0.222	0.947	0.625	0.613 ($p = 0.00039$)	0.531 ($p = 0.43$)	0.177 ($p = 6.6 \times 10^{-5}$)
CADD phred	27.500	21.200	0.848	0.628	0.648 ($p = 2.8 \times 10^{-6}$)	0.543 ($p = 0.27$)	0.181 ($p = 3.7 \times 10^{-5}$)

CDKN2A is the exception. Patient-weighted, KRAS 99.9%, TP53 98.2% and SMAD4 98.3% of missense carriers hold an AlphaMissense likely-pathogenic allele, but only 55.4% of CDKN2A missense carriers do. One candidate explanation (tested and not supported in the next section) is that the scores are defined on a single canonical protein, while CDKN2A also encodes p14^{ARF} from an alternate reading frame, so an allele’s effect on that product is not scored. We treat the low scores as a limit of the predictor, not as evidence that those alleles are passengers.

How many panel calls rest on unsupported alleles? Calling an alteration *supported* if it is truncating, ClinVar pathogenic/likely pathogenic, or AlphaMissense likely-pathogenic, 99.4% of the 3,115 panel-positive patients carry at least one supported alteration (ClinVar P/LP alone: 97.8%); 19 patients carry only unsupported alterations (mostly in-frame indels and low-scoring missense). The denominator is panel-positive patients only, so this is a bound on annotation support, not on sensitivity.




Figure 10: MyVariant allele audit. A: AlphaMissense separates driver from passenger missense. B: no predictor ranks panel alleles by recurrence robustly across cohorts. C: CDKN2A missense carriers fall below the AlphaMissense threshold far more often than the other panel genes.

20 Testing the p14^{ARF} explanation of the CDKN2A exception (negative)

The allele-level section left one explanation untested: low AlphaMissense scores on CDKN2A might mark alleles that act through p14^{ARF}, the second protein encoded in an alternate reading frame. We tested it directly. From UCSC hg19 refGene exon coordinates and the hg19 sequence we rebuilt the NM_000077 (p16^{INK4a}) and NM_058195 (p14^{ARF}) coding sequences; both translate to full-length proteins (156 and 132 aa, Met start, one terminal stop) and share 206 coding bp in exon 2. Each of the 72 CDKN2A missense SNVs that match their cBioPortal p16 call (3 excluded as annotated on other transcripts) was translated in both frames (results/cdkn2a_arf_alleles.csv).

The raw association is strong. Low-AlphaMissense alleles alter p14^{ARF} in 10/19 cases versus 7/53 for the rest (Fisher OR 7.3, $p = 0.0012$), and inside the shared exon-2 region 10/12 versus 7/33 (OR 18.6, $p = 0.00027$). The most common low-scoring allele, p16 H83Y (48 patients, AlphaMissense 0.166, ClinVar P/LP), is ARF A97V.

It is explained by reading-frame geometry. In the shared region the ARF frame is offset by one base, so a change at p16 codon position 1 alters ARF in 94% of alleles and a change at position 2 in 0%. Independently, low AlphaMissense scores concentrate at p16 codon position 1 (15/24 alleles vs 4/45 at position 2, $p = 4.4 \times 10^{-6}$). Stratified by codon position the association vanishes: in the only informative stratum (position 1) ARF is altered in 6/6 high-scoring and 10/11 low-scoring alleles (Mantel-Haenszel $p = 0.76$). **Verdict:** the data cannot separate an ARF mechanism from frame geometry; the explanation is not supported, and the CDKN2A exception is best described as AlphaMissense rating many codon-position-1 substitutions (e.g. H83Y) as benign.

Independent recurrence definition. We repeated the recurrence test with residue hotspots from cancer-hotspots.org (164 single-residue hotspots in the four panel genes; statistically defined on large pan-cancer cohorts that overlap ours, so this is a check of definition, not of data). 355 of 537 panel missense alleles sit at a hotspot residue. AlphaMissense does not separate hotspot from non-hotspot residues (AUROC 0.519, $p = 0.47$), while REVEL (0.610, $p = 2.8 \times 10^{-5}$) and CADD (0.597, $p = 0.00024$) do, reproducing the AlphaMissense panel ceiling. 13 hotspot alleles (74 of 4,647 hotspot carriers) score below the likely-pathogenic threshold, led by CDKN2A H83Y.

21 Structural context of low-scoring panel alleles (RCSB PDB, negative)

A natural reason for AlphaMissense to miss damaging alleles is that it models a single chain, so damage at a partner interface might be under-weighted. We stated this as a falsifiable hypothesis — alleles below the likely-pathogenic threshold are *enriched* at partner or ligand contact residues — and tested it on four experimental complexes from the RCSB PDB (Table 6). For each residue we computed relative solvent accessibility (RSA) of the isolated chain with the Biopython Shrake–Rupley implementation (maximum ASA from Tien et al. 2013) and contact with any partner chain, DNA, or bound ligand (any heavy atom within 5 Å, Biopython NeighborSearch). Author numbering was checked residue by residue against the UniProt canonical sequence, and an allele was mapped only when its reference amino acid matched the structure: 471 of 537 panel missense alleles mapped (results/structure_panel_alleles.csv).




Figure 11: A: fraction of CDKN2A missense alleles that alter p14^{ARF}, split by p16 codon position and AlphaMissense class (shared exon-2 region). B: AlphaMissense does not separate cancerhotspots.org residues from the rest of the panel; REVEL and CADD do.

Table 6: Structures used. Identity = share of modelled residues matching UniProt.

Gene	PDB (chain)	residues	identity %	partner contacts	ligand contacts	alleles mapped
KRAS	6VJJ (A)	168	96.4	24	32	29
TP53	1TSR (B)	194	100.0	20	6	285
CDKN2A	1BI7 (B)	125	99.2	35	0	70
SMAD4	1U7F (B)	193	100.0	65	0	87

Control. Buried residues score higher, as they should: Spearman $\rho(\text{RSA}, \text{AlphaMissense}) = -0.114$ ($p = 0.013$); REVEL $\rho = -0.100$ ($p = 0.029$); CADD, which has no structural input, shows none ($\rho = 0.012$, $p = 0.8$).

Result: the hypothesis is falsified in the opposite direction. Low-scoring alleles are *depleted* at contact residues: 3/27 versus 160/444 for the rest (Fisher OR 0.22, $p = 0.0066$; protein/DNA partner contacts alone 2/27 versus 114/444, $p = 0.036$). They are also less often buried (15/27 versus 352/444, OR 0.33, $p = 0.0074$). Both effects hold jointly with gene fixed effects in a logistic model (RSA coefficient 6.20, $p = 6.4 \times 10^{-5}$; contact coefficient -2.18, $p = 0.0022$). **Verdict:** AlphaMissense’s low scores fall mostly on exposed, non-contact residues, where low damage is structurally plausible; interface blindness does not explain them. The main exception is CDKN2A H83Y, which is fully buried (RSA 0.00) in the p16 ankyrin fold yet scored benign. The CDK6–p16 structure is 3.4 Å, and the p53 interface here is DNA only (other p53 partners are not modelled), so contact calls are a lower bound.

22 Documented negatives

(1) SMAD4 predictability is at chance across architectures. (2) The SWI/SNF–CDKN2A exclusivity fails replication. (3) Gene-level permutation importance on the CDKN2A GNN is noise. (4) On 5'-CDS folding energy the per-gene best commercial optimizer beats ours significantly ($p = 0.040$; ViennaRNA audit above). Each negative is preserved in the archived result files (results/*.json).




Figure 12: A: AlphaMissense versus relative solvent accessibility for mapped panel alleles. B: low-scoring alleles are less often at partner/ligand contacts and less often buried.

23 Falsification paths

CodonOpt's expression gains are in-silico: the scorer is our own CNN. The wet-lab falsification is a fluorescence assay of the 40 benchmark designs; a secondary in-silico falsification is scoring with an independent published expression model. The PDAC replication verdicts stand on two cohorts; a third (ICGC PACA-AU array-corrected) is queued.

24 Conclusion

PANCREAS.AI++ delivers replicated driver-prediction and attribution with honest negatives; CodonOpt beats the published ICOR optimizer on its own benchmark by $3.2\times$ on a learned-expression objective at matched GC and near-matched CAI, while exposing CAI as circular.

References

- Sharp PM, Li WH. The codon adaptation index. *Nucleic Acids Res* 1987.
- Jain R, et al. ICOR: improving codon optimization with recurrent neural networks. *BMC Bioinformatics* 2023. doi:10.1186/s12859-023-05246-8.
- Cancer Genome Atlas Research Network. Integrated genomic characterization of PDAC. *Cancer Cell* 2017.
- Grimson A, et al. MicroRNA targeting specificity in mammals. *Mol Cell* 2007.

A External tools and dataset manifest

A.1 External research/data tools and science libraries (40-tool target not yet verified)

Research source or science library	Verified use and committed evidence
cBioPortal REST API	Mutation and cohort queries (results/pancan_cdkn2a.json)
PaxDb, NCBI RefSeq	Protein abundances and E. coli CDS (codon predictor)
GTEX Portal API	Normal-tissue expression (results/gtex_panel_background.json)
GtRNAdb	tRNA counts for the tAI comparison (results/tai_benchmark.json)
Open Targets GraphQL	Disease association records (results/opentargets_panel_audit.json)
CIViC and gnomAD APIs	Clinical evidence and variant frequencies (results/ audits)
MyVariant.info and UCSC REST	Allele effects and trinucleotide context (results/ audits)
RCSB PDB, UniProt, Biopython	Structural panel and numbering (results/structure_panel_audit.json)
STRING v12 REST	Driver-panel network decomposition and CDKN2A interaction neighbourhood (results/string_panel_network.json); UniProt REST re-verified panel identities (results/uniprot_panel_check.json)
Reactome Analysis Service	Pathway enrichment of the CDKN2A STRING neighbourhood - top hits RB1/E2F G1-phase cell cycle and senescence ($\text{FDR} \sim 5 \times 10^{-15}$), independently confirming the cell-cycle attribution axis (results/reactome_neighborhood.json)
PyTorch, scikit-learn, SciPy	CNN/GNN models, baselines and statistics

The full analysis also used NumPy, pandas, matplotlib, ViennaRNA, NCBI E-utilities, DepMap, the Human Protein Atlas, ClinVar, cancerhotspots.org, statsmodels, Kazusa codon frequencies and the ICOR comparator. These are grouped here rather than padded into a count. The historical 41-item inventory incorrectly counted nine cohorts as tools and counted GTEx twice. pytest, git/GitHub, TeX and hosting are infrastructure. COSMIC was probed but inaccessible and is excluded. At most 31 of the old entries survive the obvious cohort and duplicate-service exclusions, before checking which entries are data rather than tools. STRING and Reactome were added as genuinely new verified uses on 2026-09-26 (UniProt was already counted). The strict 40-tool target remains open.

A.2 Dataset manifest (source records audited; 54,988 claim withdrawn)

The prior inventory treated 10 pancreatic cohorts and 140 additional cBioPortal study IDs as separate entries. The committed pan-cancer result has 143 study IDs in total, including at least four paired TCGA lineages under different labels. The source studies must be mapped before an independent-study count can be given. The biological record inventory can be checked directly: 9,410 distinct gnomAD variant IDs; 432 distinct GTEx gene–tissue keys; 1,146 distinct CIViC EIDs; 1,333 successfully fetched MyVariant HGVS keys; and 41,319 distinct Open Targets gene–disease pairs among 42,329 CSV rows. These records do not imply 54,988 independent experiments or clinical cohorts. No cross-source exact dataset total is asserted here.

Deduplication audit. The Open Targets CSV has 42,329 rows but only 41,319 distinct Ensembl-gene–disease pairs (1,010 duplicate pairs). The pan-cancer table includes at least four matched TCGA lineage

pairs under separate cohort IDs; an API study ID does not by itself establish an independent source study. The previous 54,988 total mixes source records with cohorts and should not be read as independent datasets. The 120-record threshold is met under the accession-record convention; independent-study threshold is not yet verified.

B Multi-cohort replication meta-analysis

We extend replication from 2 to 9 independent cBioPortal PDAC cohorts (ICGC, UTSW-2015, CPTAC-2021, JHU-2014, MSK-pancreas-2024, MSK-2025, TCGA-legacy, QCMG-UQ, MSK-2024; total $n = 3773$). Per-cohort 2×2 tables feed a Mantel–Haenszel pooled estimator

$$\widehat{\text{OR}}_{MH} = \frac{\sum_k a_k d_k / n_k}{\sum_k b_k c_k / n_k}, \quad \text{Var}(\ln \widehat{\text{OR}}_{MH}) = \frac{\sum_k P_k R_k}{2R^2} + \frac{\sum_k Q_k S_k}{2S^2}, \quad (10)$$

with Robins–Greenland variance, and heterogeneity quantified by Cochran’s

$$Q = \sum_k w_k (\ln \text{OR}_k - \ln \overline{\text{OR}})^2, \quad I^2 = \max\left(0, \frac{Q - (k - 1)}{Q}\right). \quad (11)$$

Per-cohort odds ratios feed Fig. 2 (extended version in supplement).

Covariate	pooled OR	p	Q	I^2	direction	verdict
KRAS	2.71	2.8e-13	14.27	0.439	3/4	REPLICATED
TP53 pathway	3.787	<1e-16	2.9	0.0	6/6	REPLICATED
SWI SNF	0.911	3.0e-01	13.06	0.387	4/6	FALSIFIED
RTK RAS nonKRAS	1.247	8.4e-02	5.26	0.0	5/6	FALSIFIED

Table 7: Mantel–Haenszel meta-analysis over 9 cohorts. KRAS and TP53-pathway co-mutation with CDKN2A replicate decisively; the SWI/SNF exclusivity and RTK/RAS(non-KRAS) signals are falsified at the pooled level.

C Extended derivations and training protocols

C.1 Derivation of the symmetric GCN normalization

Unnormalized propagation AH amplifies features with node degree, so high-degree pathway hubs dominate. With $\tilde{A} = A + I$ and $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$, the symmetric normalization $\tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2}$ bounds the spectral radius by 1: for any x ,

$$x^\top \tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2} x = \sum_{(i,j) \in \mathcal{E}} \left(\frac{x_i}{\sqrt{d_i}} + \frac{x_j}{\sqrt{d_j}} \right)^2 / 2 \leq \sum_i x_i^2, \quad (12)$$

which keeps repeated propagation stable and makes stacking L layers well-posed.

C.2 Gradient of the logistic head under BCE

With $\hat{y} = \sigma(z)$, $z = w^\top h + b$,

$$\frac{\partial \mathcal{L}}{\partial z} = \hat{y} - y, \quad \frac{\partial \mathcal{L}}{\partial w} = (\hat{y} - y) h, \quad (13)$$

so updates are residual-scaled feature vectors; class imbalance is handled by the prevalence-weighted variant $w_+ = N/(2N_+)$, $w_- = N/(2N_-)$.

C.3 Fisher exact p from the hypergeometric law

Conditioning on margins of the 2×2 table,

$$P(X = n_{11}) = \frac{\binom{n_{1\cdot}}{n_{11}} \binom{n_{\cdot 2}}{n_{\cdot 1} - n_{11}}}{\binom{N}{n_{\cdot 1}}}, \quad (14)$$

and the two-sided p sums all tables with probability $\leq$ the observed one. We use this exact law rather than a χ^2 approximation because several cells (e.g. SWI/SNF $\times$ CDKN2A in TCGA) contain zeros.

C.4 Permutation-importance variance estimate

For R repeats of permutation of covariate set S ,

$$\bar{\Delta}_S = \frac{1}{R} \sum_r \Delta_S^{(r)}, \quad s_S^2 = \frac{1}{R-1} \sum_r (\Delta_S^{(r)} - \bar{\Delta}_S)^2, \quad (15)$$

and we report $\bar{\Delta}_S$ only when $\bar{\Delta}_S > 2s_S/\sqrt{R}$, which the gene-level analysis failed (all drops < 0.004 AUC) and the pathway-level analysis passed for RTK/RAS and SWI/SNF.

C.5 CAI properties and the circularity theorem

Theorem (CAI circularity). Let $c^*(a) = \arg \max_{c \sim a} f(c)$ per amino acid a and let s^* map every residue to c^* . Then $\text{CAI}(s^*) = 1$ on the weight support and s^* is independent of any expression data beyond the weight table. *Proof.* Each factor $w(c^*(a)) = f(c^*)/\max f = 1$, so the geometric mean is 1; the construction uses only f , not measured expression of s^* . Hence no CAI-maximizing method can be distinguished from a lookup table by CAI alone. $\square$ On the benchmark, HFC instantiates s^* at CAI 0.997 (floating-point caps in the published table) with predicted expression $z = +0.60$, below ICOR's $+0.68$ and far below CodonOpt's $+2.20$.

C.6 ExpressionCNN training protocol

Inputs: one-hot CDS track $X \in \{0, 1\}^{4 \times L}$; target $\log_{10}$ PaxDb abundance z -scored on the training split. Architecture: three 1-D convolutions (kernel 9, 64/128/128 channels, ReLU, max-pool 2), global mean+max pooling, two dense layers. Optimization:

$$\theta^* = \arg \min_{\theta} \frac{1}{N} \sum_i (\hat{z}_i - z_i)^2 + \lambda \|\theta\|_2^2, \quad \lambda = 10^{-4}, \quad (16)$$

Adam, lr 10^{-3} , batch 256, early stop on held-out-gene Spearman. Final held-out $\rho_s = 0.61$.

C.7 Constrained hill-climb optimizer

Objective (6) is optimized by steepest-ascent codon swaps:

1. initialize s at the wild-type CDS;
2. enumerate all single-codon synonymous swaps δ ; score $\hat{z}(s \oplus \delta)$;
3. apply the best swap satisfying $\text{GC} \in [0.50, 0.56]$ and $\text{CAI} \geq 0.80$;
4. stop when no feasible improving swap exists.

Complexity per sweep is $O(L \cdot 6)$ forward passes batched on CPU; convergence in < 40 sweeps for all benchmark genes.

C.8 GC content as a matching constraint

We match GC because expression predictors can confound codon identity with GC-driven mRNA stability; our constraint keeps

$$|\text{GC}(s) - \text{GC}(s_{\text{ICOR}})| \leq 0.03 \quad (17)$$

per gene (achieved mean 0.538 vs 0.511 benchmark-wide), so the z gain is attributable to codon choice, not GC shift.

D Supplementary tables

D.1 Table S1: per-gene head-to-head on the ICOR 40-gene benchmark

Gene	CAI			pred expr z			ours	
	WT	ICOR	ours	WT	ICOR	ours	GC	Δz vs ICOR
BIRC5	0.597	0.863	0.691	-0.19	+0.67	+2.10	0.505	+1.43
BRAF1	0.528	0.841	0.901	-0.67	+0.15	+2.27	0.562	+2.12
CAV1	0.661	0.892	0.731	+0.20	+0.12	+2.22	0.461	+2.10
CD80	0.554	0.875	0.757	-0.36	+0.78	+2.15	0.510	+1.37
CDK1	0.523	0.891	0.755	-0.43	+0.78	+2.48	0.489	+1.70
CEBPZ	0.573	0.904	0.924	-0.37	+1.25	+2.45	0.490	+1.20
CLN3	0.598	0.847	0.852	-0.48	-0.03	+1.65	0.581	+1.68
CREB1	0.534	0.862	0.780	-0.02	+1.42	+2.62	0.578	+1.20
CSNK1A1	0.543	0.893	0.802	-0.11	+1.17	+2.60	0.496	+1.43
FGFR4	0.581	0.863	0.897	-0.27	+0.83	+2.06	0.614	+1.23
GSK3B	0.530	0.860	0.838	-0.62	+0.72	+2.26	0.550	+1.54
JUN	0.629	0.851	0.743	-0.55	+0.60	+2.60	0.571	+2.00
KIF11	0.550	0.873	0.918	-0.12	+1.50	+2.51	0.504	+1.01
LAMP1	0.584	0.837	0.814	-0.52	+0.21	+2.18	0.548	+1.97
LEMD3	0.531	0.852	0.920	-0.28	+1.50	+2.45	0.595	+0.95
MAPK1	0.573	0.892	0.786	-0.10	+0.58	+2.42	0.497	+1.84
MAPKAPK5	0.550	0.906	0.843	+0.20	+1.05	+2.53	0.518	+1.48
NGFR	0.642	0.861	0.826	-0.20	+0.61	+2.33	0.609	+1.72
NOC2L	0.590	0.866	0.906	+0.08	+0.85	+2.09	0.562	+1.24
OPRM1	0.584	0.862	0.807	-0.92	-0.32	+1.30	0.508	+1.62
PDCD11	0.561	0.869	0.956	+0.19	+1.01	+2.04	0.544	+1.03
PLK1	0.573	0.875	0.865	+0.03	+0.79	+2.31	0.549	+1.52
RPS6KB1	0.516	0.889	0.851	-0.74	+1.10	+2.37	0.548	+1.27
SMARCD1	0.562	0.884	0.843	-0.13	+0.70	+2.46	0.574	+1.76
TAP1	0.570	0.853	0.897	-0.65	+0.01	+1.81	0.612	+1.80
TAS2R10	0.517	0.869	0.815	-1.29	-0.50	+1.57	0.458	+2.07
akt1	0.587	0.895	0.842	-0.17	+0.92	+2.39	0.513	+1.47
emg1	0.531	0.866	0.724	-0.16	+1.28	+2.41	0.520	+1.13
falvac-1	0.592	0.888	0.778	-0.15	+0.80	+2.51	0.445	+1.71
flt1	0.557	0.864	0.935	-1.01	-0.11	+1.24	0.521	+1.35
hpdf	0.664	0.899	0.673	-0.33	+0.19	+1.93	0.592	+1.74
lck	0.630	0.885	0.853	-0.10	+0.53	+1.96	0.532	+1.43
mmp13	0.665	0.848	0.919	+0.05	+0.54	+1.91	0.624	+1.37
npr1	0.612	0.871	0.916	-0.33	+0.26	+1.56	0.593	+1.30
pa	0.517	0.923	0.716	-0.40	+1.14	+2.20	0.457	+1.06
pak1	0.559	0.886	0.850	+0.35	+1.00	+2.87	0.526	+1.87
pea	0.743	0.882	0.890	-0.11	+0.84	+2.09	0.617	+1.25
pim1	0.585	0.890	0.784	-0.39	+0.69	+2.51	0.546	+1.82
ptp4a3	0.676	0.897	0.696	+0.21	+1.16	+2.44	0.508	+1.28
ubtf	0.577	0.893	0.877	+0.09	+0.57	+2.17	0.485	+1.60

ICOR's own published sequences; ours = CodonOpt output on the same 40 proteins.

D.2 Table S2: E. coli codon-optimization panel (all genes)

Locus	gene	len (aa)	WT CAI	ours CAI	WT z	ours z
b4621	yjbS	67	0.539	0.617	-0.86	+1.84
b1847	yebF	118	0.685	0.619	+0.59	+2.47
b3642	pyrE	213	0.751	0.623	+0.76	+2.55
b4278	insG	442	0.606	0.675	-0.58	+2.44
b2026	hisI	203	0.744	0.645	+0.63	+2.50
b0190	yaeQ	181	0.750	0.621	-0.06	+2.26
b1244	oppB	306	0.649	0.615	-0.49	+2.08
b2469	narQ	566	0.687	0.732	-0.74	+2.46
b2333	yfcP	179	0.660	0.635	-0.41	+2.42
b1938	fliF	552	0.723	0.753	+0.23	+2.85

Greedy-CAI reaches CAI = 1.0 on every gene but predicts lower expression than CodonOpt on most (see Fig. 4).

D.3 Table S3: co-mutation 2×2 tables, all cohorts

Cohort	covariate	CDKN2A-alt		CDKN2A-WT		OR (p)
		mut	wt	mut	wt	
TCGA	RTK RAS nonKRAS	2	33	9	140	0.94 (1.0e+00)
TCGA	KRAS	33	2	84	65	12.77 (8.9e-06)
TCGA	SWI SNF	0	35	15	134	0.00 (7.9e-02)
TCGA	DNA REPAIR	0	35	4	145	0.00 (1.0e+00)
TCGA	TP53 pathway	31	4	83	66	6.16 (2.0e-04)
TCGA	TGF BETA	9	26	37	112	1.05 (1.0e+00)
paad qcmg uq 2016	KRAS	65	6	279	106	4.12 (2.8e-04)
paad qcmg uq 2016	TP53 pathway	60	11	207	178	4.69 (7.2e-07)
paad qcmg uq 2016	SWI SNF	13	58	38	347	2.05 (6.2e-02)
paad qcmg uq 2016	RTK RAS nonKRAS	1	70	19	366	0.28 (3.4e-01)
pdac msk 2024	KRAS	525	21	1663	127	1.91 (6.5e-03)
pdac msk 2024	TP53 pathway	500	46	1327	463	3.79 (2.0e-20)
pdac msk 2024	SWI SNF	61	485	255	1535	0.76 (7.4e-02)
pdac msk 2024	RTK RAS nonKRAS	30	516	82	1708	1.21 (4.2e-01)

Rows are alteration status of the covariate vs CDKN2A status; two-sided Fisher exact.

E Supplementary table: per-gene multi-objective results

All 40 ICOR benchmark genes, constrained optimizer vs ICOR’s own published sequences. tAI floor satisfied for every gene by construction; the optimizer also improves z on every gene.

Gene	tAI (ours)	tAI (ICOR)	z (ours)	z (ICOR)	CAI (ours)	GC (ours)
BIRC5	0.3486	0.3471	2.12	0.67	0.716	0.513
BRAF1	0.3589	0.3285	2.60	0.15	0.697	0.554
CAV1	0.3483	0.3477	2.26	0.12	0.677	0.451
CD80	0.3351	0.3324	2.34	0.78	0.659	0.499
CDK1	0.3703	0.3675	2.52	0.78	0.654	0.488
CEBPZ	0.4069	0.3860	2.81	1.25	0.727	0.490
CLN3	0.3357	0.3102	2.05	-0.03	0.670	0.573
CREB1	0.3369	0.3248	2.78	1.42	0.645	0.542
CSNK1A1	0.3832	0.3829	2.71	1.17	0.683	0.493
FGFR4	0.3610	0.3270	2.30	0.83	0.680	0.588
GSK3B	0.3419	0.3143	2.59	0.71	0.666	0.542
JUN	0.3328	0.3265	2.68	0.60	0.649	0.570
KIF11	0.3962	0.3672	2.82	1.50	0.735	0.492
LAMP1	0.3564	0.3316	2.40	0.21	0.655	0.532
LEMD3	0.3765	0.3328	2.80	1.50	0.695	0.572
MAPK1	0.3669	0.3659	2.58	0.58	0.670	0.494
MAPKAPK5	0.3686	0.3653	2.67	1.05	0.707	0.516
NGFR	0.3341	0.2858	2.65	0.61	0.645	0.556
NOC2L	0.3730	0.3470	2.40	0.85	0.709	0.548
OPRM1	0.3368	0.3220	1.62	-0.32	0.658	0.515
PDCD11	0.3983	0.3611	2.49	1.01	0.709	0.517
PLK1	0.3674	0.3532	2.48	0.79	0.683	0.549
RPS6KB1	0.3564	0.3549	2.46	1.10	0.832	0.540
SMARCD1	0.3335	0.3321	2.47	0.70	0.815	0.569
TAP1	0.3623	0.3391	2.05	0.01	0.679	0.580
TAS2R10	0.3606	0.3588	1.65	-0.50	0.764	0.447
akt1	0.3691	0.3684	2.58	0.92	0.820	0.510
emg1	0.3602	0.3592	2.44	1.28	0.635	0.511
falvac-1	0.3583	0.3580	2.62	0.80	0.664	0.466
flt1	0.3819	0.3468	1.60	-0.11	0.715	0.523
hpdf	0.3194	0.3183	1.87	0.19	0.644	0.570
lck	0.3367	0.3353	2.09	0.53	0.825	0.534
mmpl3	0.3691	0.3322	2.34	0.54	0.682	0.591
npr1	0.3863	0.3502	2.23	0.26	0.716	0.557
pa	0.3818	0.3799	2.19	1.14	0.721	0.461
pak1	0.3458	0.3427	2.97	1.00	0.830	0.519
pea	0.3505	0.3316	2.35	0.84	0.697	0.584
pim1	0.3413	0.3409	2.52	0.69	0.658	0.527
ptp4a3	0.3477	0.3387	2.51	1.16	0.650	0.481
ubtf	0.3966	0.3897	2.42	0.57	0.690	0.495

F Reproduction

All figures and tables regenerate from `results/*.json` via `paper/make_figures.py`; models retrain via `scripts/*.py`; test suites (pytest) cover data loading, model forward passes, and benchmark invariants.

G Appendix: per-gene ICOR head-to-head under the source-pinned check-point

All 40 benchmark genes, our CNN v2 design vs ICOR, scored under the source-pinned expression check-point (results/icor_rescore.json; paired Wilcoxon $p = 1.8 \times 10^{-12}$, $n = 40$). Ours beats ICOR’s predicted expression z on 40/40 genes at matched-or-better GC.

Gene	ours (CNN v2)			ICOR		
	CAI	GC	z	CAI	GC	z
BIRC5	0.686	0.529	+1.12	0.863	0.512	+0.17
BRAF1	0.888	0.562	+1.24	0.841	0.523	-0.14
CAV1	0.679	0.483	+1.34	0.892	0.434	+0.15
CD80	0.759	0.520	+1.12	0.875	0.505	+0.19
CDK1	0.748	0.502	+1.56	0.891	0.473	+0.57
CEBPZ	0.928	0.491	+1.61	0.904	0.457	+1.06
CLN3	0.811	0.579	+0.67	0.847	0.566	-0.13
CREB1	0.812	0.590	+1.77	0.862	0.555	+0.87
CSNK1A1	0.765	0.501	+1.32	0.893	0.465	+0.47
FGFR4	0.894	0.613	+1.23	0.863	0.570	+0.49
GSK3B	0.821	0.556	+1.25	0.860	0.528	+0.47
JUN	0.763	0.579	+1.59	0.851	0.559	+0.11
KIF11	0.932	0.507	+1.69	0.873	0.468	+0.83
LAMP1	0.803	0.556	+1.18	0.837	0.528	-0.17
LEMD3	0.910	0.596	+2.12	0.852	0.543	+1.42
MAPK1	0.795	0.506	+1.51	0.892	0.463	+0.20
MAPKAPK5	0.819	0.506	+1.59	0.906	0.489	+0.55
NGFR	0.810	0.619	+1.48	0.861	0.586	+0.41
NOC2L	0.877	0.563	+1.28	0.866	0.530	+0.55
OPRM1	0.793	0.515	+0.67	0.862	0.498	-0.44
PDCD11	0.951	0.546	+1.42	0.869	0.504	+0.40
PLK1	0.874	0.557	+1.36	0.875	0.504	+0.48
RPS6KB1	0.841	0.543	+1.73	0.889	0.504	+0.85
SMARCD1	0.824	0.580	+1.53	0.884	0.535	+0.24
TAP1	0.887	0.616	+0.91	0.853	0.575	-0.17
TAS2R10	0.782	0.473	+0.65	0.869	0.418	-0.59
akt1	0.835	0.510	+1.46	0.895	0.492	+0.56
emg1	0.723	0.549	+1.47	0.866	0.516	+0.65
falvac-1	0.780	0.445	+1.53	0.888	0.431	+0.08
flt1	0.934	0.523	+0.72	0.864	0.482	-0.30
hpdf	0.685	0.595	+1.13	0.899	0.615	+0.17
lck	0.824	0.542	+1.09	0.885	0.516	+0.26
mmp13	0.919	0.624	+1.34	0.848	0.565	+0.43
npr1	0.915	0.593	+0.95	0.871	0.542	-0.14
pa	0.720	0.478	+1.41	0.923	0.433	+0.53
pak1	0.838	0.536	+1.92	0.886	0.490	+0.81
pea	0.878	0.621	+1.44	0.882	0.568	+0.35
pim1	0.765	0.548	+1.12	0.890	0.521	+0.10
ptp4a3	0.678	0.558	+1.50	0.897	0.522	+0.57
ubtf	0.871	0.494	+1.41	0.893	0.455	+0.36

H Pan-cancer generalization: CDKN2A landscape across 143 API study IDs

To test whether the PDAC-core CDKN2A finding is pancreas-specific, we fetched a four-gene driver panel (KRAS, TP53, CDKN2A, SMAD4) for 143 cBioPortal API study IDs (size-stratified across small and large studies, $n \geq 25$; 140 not previously used; not 159 independent studies). Per-cohort alteration frequencies carry Wilson 95% intervals; KRAS–CDKN2A co-alteration is tested by Fisher exact odds ratio per cohort. CDKN2A prevalence is low in the median cancer cohort (median 1.4%, IQR 0–4.9% across 143

Study	n	CDKN2A freq	95% CI	KRAS-CDKN2A OR
csc ranson 2022	25	0.800	[0.609,0.911]	-
csc dfarber 2015	29	0.448	[0.284,0.625]	0.0
sarcoma dfci genie 2026	48	0.333	[0.217,0.475]	2.067
csc ucsf 2021	83	0.265	[0.182,0.369]	0.0
pancreas cptac gdc	183	0.213	[0.160,0.278]	0.914
paad cptac 2021	140	0.207	[0.148,0.282]	0.375
skcm vanderbilt mskcc 2015	66	0.197	[0.119,0.308]	-
skcm broad	121	0.190	[0.130,0.269]	-
lusc tcga pub	178	0.174	[0.126,0.237]	0.0
lusc cptac gdc	110	0.164	[0.106,0.244]	1.294
ohnca cptac gdc	172	0.163	[0.115,0.225]	-
biliary tract adc targets msk 2026	33	0.151	[0.067,0.309]	0.0

Table 9: Top cohorts by CDKN2A alteration frequency (full per-cohort table in `results/pancan_cdkn2a.json`).

API study IDs) but sharply cancer-type-specific: the top of the ranking is occupied by cutaneous squamous cell carcinoma cohorts (27–80%: `csc_ucsf_2021`, `csc_dfarber_2015`, `csc_ranson_2022`), consistent with UV-driven CDKN2A loss as a known cSCC driver, and the PDAC cohorts sit in the high-prevalence band (10–21% for CPTAC/MSK; `paad_cptac_2021` 20.7%, `pancreas_cptac_gdc` 21.3%), replicating the item’s core finding across multiple API study IDs (independence not fully established). Two honest data notes: `paad_icgc` shows only 2.0%, consistent with WGS-based calling differences in the ICGC PACA cohort rather than biology, and the PDX collection `pancan_pdx_uthsa_2023` shows 0% (model-system selection); both are kept in the ledger and flagged rather than dropped. Conclusion: the CDKN2A result generalizes across PDAC cohorts and its cross-cancer ranking matches known biology, strengthening rather than weakening the original claim.