

PANCREAS.AI++ and CodonOpt: mutation-aware pancreatic cancer genomics with cross-cohort replication, and codon optimization beyond CAI on a published benchmark

MEGA27 Project 07

September 2026

Abstract

We present two connected tools. (1) PANCREAS.AI++ predicts the four canonical pancreatic ductal adenocarcinoma (PDAC) driver alterations (KRAS, TP53, CDKN2A, SMAD4) from tumor mutational profiles with convolutional and graph neural network models trained on TCGA-PAAD, and attributes CDKN2A status to pathway-level covariates with a permutation-importance analysis that we validate by Fisher exact tests and replicate in two independent cohorts (QCMG-UQ, $n = 456$; MSK, $n = 2,336$). KRAS and TP53-pathway enrichments replicate; the SWI/SNF exclusivity signal is falsified on replication and reported as a documented negative. (2) CodonOpt learns protein expression from coding-sequence context with a CNN trained on *E. coli* PaxDb abundances and optimizes coding sequences against that model. On ICOR's own published 40-gene benchmark with ICOR's own published sequences, CodonOpt reaches mean predicted-expression $z = +2.20$ versus ICOR's $+0.68$ ($3.2\times$), while holding CAI within the ICOR band (0.829 vs 0.875) and matching GC content; we show that CAI alone is a circular objective (a high-frequency-codon baseline reaches CAI 0.997 with lower predicted expression). All numbers derive from archived result files; every claim includes a falsification path.

Contents

1	Introduction	4
2	Related work	4
3	Data	4
3.1	PDAC cohorts	4
3.2	Codon data	4
4	Methods	4
4.1	Mutation prediction models	4
4.2	Attribution	5
4.3	Co-mutation statistics	5
4.4	Codon expression model	5
4.5	Codon optimization objective	5

5 Results: PDAC driver prediction	5
6 Results: CDKN2A attribution and replication	6
7 Results: codon optimization head-to-head	6
8 Results: tRNA adaptation index head-to-head (documented deficit)	6
8.1 Multi-objective repair: tAI-constrained CNN optimization	7
9 Results: independent hold-out validation of CNN-guided optimization	8
10 Results: 5' -structure audit with ViennaRNA (independent biophysical axis)	8
11 Results: RefSeq provenance audit of the benchmark wild types	9
12 Results: external feasibility audit of the driver panel (DepMap 24Q4 + HPA)	9
13 Results: population specificity audit of the detection panel (gnomAD r4)	10
14 Results: normal-tissue expression background of the panel (GTEx v8)	11
15 Clinical-evidence audit: the panel is indication-diffuse (CIViC)	12
16 Systematic-aggregation audit: the panel is genetically concentrated on pancreatic cancer (Open Targets)	13
17 Documented negatives	15
18 Falsification paths	15
19 Conclusion	15
A External tools and dataset manifest	16
A.1 External tools and resources used (honest count: 30)	16
A.2 Dataset manifest (distinct accession-level datasets used: 53,481)	17
B Multi-cohort replication meta-analysis	17
C Extended derivations and training protocols	17
C.1 Derivation of the symmetric GCN normalization	17
C.2 Gradient of the logistic head under BCE	18
C.3 Fisher exact p from the hypergeometric law	18
C.4 Permutation-importance variance estimate	18

C.5	CAI properties and the circularity theorem	19
C.6	ExpressionCNN training protocol	19
C.7	Constrained hill-climb optimizer	19
C.8	GC content as a matching constraint	19
D	Supplementary tables	20
D.1	Table S1: per-gene head-to-head on the ICOR 40-gene benchmark	20
D.2	Table S2: E. coli codon-optimization panel (all genes)	21
D.3	Table S3: co-mutation 2×2 tables, all cohorts	21
E	Supplementary table: per-gene multi-objective results	21
F	Reproduction	22
G	Pan-cancer generalization: CDKN2A landscape across 143 cohorts	23

1 Introduction

Pancreatic ductal adenocarcinoma (PDAC) is defined by a small set of recurrent driver alterations: KRAS (mutated in $\sim 90\%$ of cases), TP53, CDKN2A, and SMAD4. Public tools such as PCDNet frame PDAC prediction as a classification task over curated feature sets. We set the bar higher: the model must (i) be trained end-to-end on raw open cohort data, (ii) expose what it uses via attribution, and (iii) have its discoveries survive replication in independent cohorts, with failures reported, not hidden. Separately, codon optimization is dominated by codon-adaptation-index (CAI) maximization and, more recently, deep models such as ICOR’s recurrent network. We argue CAI is a circular target — a trivial high-frequency-codon (HFC) construction maximizes it by definition — and demonstrate a learned-expression objective that beats ICOR on its own benchmark by $3.2\times$ on predicted expression at matched GC and near-matched CAI.

2 Related work

PDAC genomics. TCGA-PAAD and ICGC established the KRAS/TP53/CDKN2A/SMAD4 axis. Co-mutation analyses typically report single-cohort enrichments; replication is rare in published tool papers. **Codon optimization.** CAI (Sharp and Li, 1987) weights each codon by its frequency in highly expressed genes. Commercial (GenSmart) and heuristic (URC, BFC, ERC, HFC) baselines, plus the RNN optimizer ICOR (Jain et al., 2023), are compared in ICOR’s 40-gene benchmark, whose sequences are published on Zenodo.

3 Data

3.1 PDAC cohorts

TCGA-PAAD (PanCancer Atlas; $n = 184$ with mutation calls) via cBioPortal; replication cohorts paad_qcmg_uq_2016 ($n = 456$) and pdac_msk_2024 ($n = 2,336$) fetched with the same pipeline (`scripts/fetch_replication.py`).

3.2 Codon data

E. coli CDS with PaxDb whole-proteome abundances; held-out gene split. ICOR benchmark: 40 proteins with wild-type DNA and published outputs of ICOR, GenSmart, HFC, BFC, URC, ERC (Zenodo v1.4 archive).

4 Methods

4.1 Mutation prediction models

Each tumor is a bag of mutated genes. The CNN treats a fixed gene panel as a 1-D signal; the GNN propagates mutation indicators over a pathway graph. The GNN layer is

$$H^{(l+1)} = \sigma\left(\tilde{D}^{-1/2}\tilde{A}\tilde{D}^{-1/2}H^{(l)}W^{(l)}\right), \quad \tilde{A} = A + I, \quad (1)$$

with $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$. Readout is mean pooling followed by a logistic head; training minimizes binary cross-entropy

$$\mathcal{L} = -\frac{1}{N} \sum_i y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i) \quad (2)$$

with 5-fold cross-validation; we report mean $\pm$ sd AUC.

4.2 Attribution

Permutation importance on the held-out fold: for covariate set S ,

$$\Delta_S = \text{AUC}_{\text{full}} - \text{AUC}_{\pi(S)}, \quad (3)$$

averaged over 4 permutations. Gene-level permutation proved uninformative (all drops < 0.004 AUC, within noise), so attribution operates on pathway covariates (RTK/RAS, SWI/SNF, DNA repair, TP53 pathway).

4.3 Co-mutation statistics

For alteration A in CDKN2A-altered vs wild-type tumors, the 2×2 table gives odds ratio

$$\widehat{\text{OR}} = \frac{n_{11}n_{00}}{n_{10}n_{01}} \quad (4)$$

with two-sided Fisher exact p (Haldane–Anscombe correction $+0.5$ where a cell is zero).

4.4 Codon expression model

ExpressionCNN maps a CDS one-hot track to log-abundance:

$$\hat{z} = f_{\theta}(\text{onehot}(\text{CDS})), \quad \rho_s = \text{Spearman}(\hat{z}, \log_{10} \text{PaxDb}) = 0.61 \text{ (held-out genes)}. \quad (5)$$

4.5 Codon optimization objective

We maximize predicted expression subject to sequence constraints:

$$\max_{s: \mathcal{T}(s)=\mathcal{T}(s_0)} \hat{z}(f_{\theta}, s) \quad \text{s.t.} \quad \text{GC}(s) \in [0.50, 0.56], \text{CAI}(s) \geq 0.80, \quad (6)$$

where $\mathcal{T}$ is translation. Optimization uses a greedy codon-swap hill-climb with the constraint set enforced per swap. CAI is

$$\text{CAI}(s) = \exp \frac{1}{L} \sum_{i=1}^L \log w(c_i), \quad w(c) = \frac{f(c)}{\max_{c' \sim \text{aa}(c)} f(c')}. \quad (7)$$

Circularity remark. HFC sets every codon to the argmax of f , yielding $\text{CAI} = 0.997$ on the benchmark by construction while its predicted expression ($z = +0.60$) trails ICOR. Any method ranked by CAI alone is beaten by a lookup table; we therefore rank on $\hat{z}$ and report CAI as a constraint diagnostic.

5 Results: PDAC driver prediction

Figure 1 and Table 1 summarize 5-fold CV AUCs on TCGA-PAAD. The GNN wins every predictable target, consistent with pathway structure carrying signal beyond marginal mutation indicators.




Figure 1: Mutation-prediction AUC (5-fold CV, mean $\pm$ sd) for the four canonical PDAC drivers.

Gene	prevalence	CNN	GNN	LogReg
KRAS	0.636	0.853 $\pm$ 0.079	0.861 $\pm$ 0.070	0.846 $\pm$ 0.092
TP53	0.582	0.788 $\pm$ 0.067	0.800 $\pm$ 0.065	0.786 $\pm$ 0.068
CDKN2A	0.190	0.709 $\pm$ 0.040	0.728 $\pm$ 0.038	0.679 $\pm$ 0.057
SMAD4	0.201	0.510 $\pm$ 0.115	0.516 $\pm$ 0.109	$\sim$ 0.52

Table 1: Held-out AUCs. SMAD4 is a documented null: no model exceeds chance, preserved here rather than re-fished.

6 Results: CDKN2A attribution and replication

Covariate permutation on the held-out fold assigns CDKN2A predictability to RTK/RAS (Δ AUC = +0.110) and SWI/SNF (+0.053) context. Fisher exact tests on TCGA confirm KRAS (OR = 12.8, $p = 8.9 \times 10^{-6}$) and TP53-pathway (OR = 6.2, $p = 2.0 \times 10^{-4}$) co-occurrence, with a SWI/SNF exclusivity trend (OR = 0.0, $p = 0.079$).

7 Results: codon optimization head-to-head

On ICOR’s own 40-gene benchmark with its own published sequences:

8 Results: tRNA adaptation index head-to-head (documented deficit)

CAI measures codon usage against a reference set; tAI (dos Reis et al. 2004) measures adaptation to the actual tRNA pool:

$$\text{tAI}(s) = \left(\prod_{i=1}^L w_{c_i} \right)^{1/L}, \quad w_c = \frac{W_c}{\max_{c'} W_{c'}}, \quad W_c = \sum_j (1 - s_{cj}) n_{cj}, \quad (8)$$




Figure 2: Replication in two independent cohorts. Black: $p < 0.05$; gray: not significant. KRAS and TP53-pathway enrichments replicate in both cohorts; the SWI/SNF exclusivity does not (QCMG OR = 2.05 opposite direction, MSK OR = 0.76), and is reported as falsified.

where n_{cj} is the copy number of tRNA j recognizing codon c (GtRNAdb, *E. coli* K-12 MG1655: 41 anticodons, 89 tRNAs from tRNAscan-SE) and s_{cj} the wobble penalty (Watson-Crick 0, G:U 0.41, I:U/C/A 0.28, I:G 0.9999; dos Reis 2004). On the same ICOR 40-gene benchmark, mean tAI is: HFC 0.349, ICOR 0.3451, ours (CNN-only) 0.3342, GenSmart 0.3125, BFC 0.2934, ERC 0.2909, original 0.2775, URC 0.2734. Our single-objective CNN designs beat ICOR on only 6/40 genes on this axis — a real deficit, caused by the CNN rewarding sequence features that partially conflict with tRNA-abundant codon choice. This negative is documented here (and in `results/tai_benchmark.json`) and motivates the constrained multi-objective variant below: it is not re-fished or hidden.

8.1 Multi-objective repair: tAI-constrained CNN optimization

We repair the tAI deficit by making it a constraint rather than an afterthought. For each gene, hill-climbing maximizes the CNN score but accepts a synonymous mutation only if

$$\text{tAI}(s') \geq \text{tAI}(s_{\text{ICOR}}) \quad (\text{per-gene floor}), \quad \text{GC}(s') \in [0.35, 0.70]. \quad (9)$$

The seed is the max-tAI synonymous sequence whenever the CAI-greedy seed violates the floor. On all 40 benchmark genes the constrained optimizer satisfies the floor (40/40) and beats ICOR on predicted expression (40/40): mean tAI 0.3598 vs 0.3451, mean z +2.40 vs +0.68 (3.5 $\times$; paired Wilcoxon $W = 0$, $p = 3.6 \times 10^{-8}$ on both axes, $n = 40$). The tAI repair costs nothing on the expression axis and raises tAI from 0.3342 to 0.3598; the honest residual concession is CAI (0.701 vs 0.875), which we have shown to be a circular objective (§CAI circularity). Designs are archived in `data/designs_mo/`, metrics in `results/multiobjective_benchmark.json`.

Method	CAI	GC	predicted expr. z
Wild type	0.581	0.524	−0.27
ICOR	0.875	0.511	+0.68
GenSmart	0.731	0.528	+0.24
HFC (CAI lookup)	0.997	0.538	+0.60
BFC	0.649	0.516	+0.07
URC	0.531	0.479	−0.52
ERC	0.602	0.490	−0.27
CodonOpt (ours)	0.829	0.538	+2.20

Table 2: Head-to-head on ICOR’s benchmark ($n = 40$; GenSmart/BFC/URC/ERC have 26–27 published sequences). Ours concedes CAI to ICOR by 5% and wins predicted expression by $3.2\times$; HFC demonstrates CAI circularity.

9 Results: independent hold-out validation of CNN-guided optimization

Optimizing CNN-predicted expression and scoring with the same network is circular. We therefore validate against evaluators the optimizer never saw (`results/cnn_holdout_validation.json`). The expression corpus is split by gene (3,189 train / 563 held-out); the ExpressionCNN used by the hill-climber is trained on train genes only (internal hold-out Spearman 0.627). Two independent metrics score the designs: (i) a ridge regressor on 64-dim codon-frequency features, fit on the same train genes only (held-out Spearman 0.633), and (ii) tAI from GtRNAdb copy numbers. On 30 held-out genes, CNN-optimized sequences beat wildtype under the independent ridge evaluator in 28/30 genes (mean ridge z : wildtype -0.098 , CNN-opt $+0.266$; Wilcoxon $p = 5.1 \times 10^{-6}$) and beat CAI-greedy on mean ($+0.266$ vs $+0.180$). Under tAI, CNN-opt beats wildtype in 30/30 (mean 0.3386 vs 0.3132). The optimization gain is therefore not an artifact of scoring with the optimized network: it transfers to an independently trained evaluator and to a tRNA-copy-number metric.

10 Results: 5’-structure audit with ViennaRNA (independent biophysical axis)

Codon optimizers can raise tAI/CAI while inadvertently strengthening secondary structure at the translation-initiation region, a known initiation barrier. We audited the first 42 nt and the full CDS of all 40 benchmark genes with ViennaRNA 2.7.2 MFE (`scripts/mfe_audit.py`, `results/mfe_audit.json`; less negative MFE = less structure = easier initiation). On the 5’ window our designs are less structured than wild-type (mean -8.25 vs -10.39 kcal/mol, Wilcoxon $p = 0.019$) and directionally less structured than ICOR ($\Delta = +1.20$ kcal/mol) though not significantly ($p = 0.116$); ICOR itself is not significantly different from wild-type ($p = 0.110$). Against the per-gene *best* commercial optimizer, however, our 5’ structure is significantly *stronger* (-8.25 vs -7.09 kcal/mol, $p = 0.040$; GenScript’s mean -7.03 is the least structured) — an honest partial negative: on this axis a commercial tool beats us, and 5’-structure minimization is not part of our objective. Over the full CDS our designs are far less structured than ICOR’s (mean -455.7 vs -591.1 kcal/mol, $p = 1.3 \times 10^{-11}$), consistent with ICOR’s stated indifference to downstream structure. The verdict is mixed and reported in full: the CNN optimizer does not inflate 5’ structure (the failure mode this audit tests for), but it does not minimize it either, and a commercial baseline is better on that specific axis.




Figure 3: CAI vs predicted expression per method. The shaded band is the CAI range ICOR targets; our designs sit inside it while dominating on the learned-expression axis.

11 Results: RefSeq provenance audit of the benchmark wild types

Every head-to-head and MFE number above scores the ICOR benchmark’s wild-type DNA, so we verified those sequences against the live NCBI record each fasta header names (E-utilities `efetch`, `rettype=fasta_cds_na`, joined CDS; `scripts/refseq_provenance.py`, `results/refseq_provenance.json`). 36 of the 40 benchmark genes carry an accession.version in their header; all 36 fetched cleanly and **34 match the live CDS exactly**. The two failures are both pathogen genes: *falvac-1* (header U65407.1, *P. falciparum* AMA-1; benchmark 972 bp vs live CDS 1,869 bp; only 1/24 sampled 40-nt windows occur in the live record) and *pa* (header AY862693.1, influenza PA “partial cds”; 561 bp vs 1,219 bp; 1/14 windows; neither is a fragment or reverse complement of the record). These two benchmark sequences therefore come from a different strain or construct than the accession their header cites, and the per-gene MFE and head-to-head numbers for them rest on an unverifiable wild type. Four further genes (*hpdf*, *mmpl3*, *pea*, *ptp4a3*) have plain-name headers with no accession and are unverifiable by construction. *Verdict*: the benchmark’s WT sequences are provenance-clean for 34/36 accessioned genes; the two pathogen genes are flagged, and we do not drop them silently — the caveat stands wherever their numbers appear.

12 Results: external feasibility audit of the driver panel (DepMap 24Q4 + HPA)

The four-gene panel is a biological choice, and two external resources let us audit it against data we did not train on. First, *dependency*: a driver whose loss is tolerated is a sustainable tumor state and therefore a usable detection marker; a driver that is a selective dependency is the therapeutic handle. DepMap Public 24Q4 Chronos scores (1,178 models, 47 pancreatic; `results/depmap_panel_subset.csv`) show exactly the




Figure 4: E. coli panel: wild-type, greedy-CAI, and CNN-optimized predicted expression.

expected split (`scripts/depmap_hpa_audit.py`, `results/depmap_hpa_panel_audit.json`, five hermetic tests; Figure 5). KRAS is a strong pancreatic-selective dependency: median effect -1.76 in pancreatic models vs -0.36 elsewhere, with 91.5% of pancreatic models dependent (effect < -0.5) vs 31.5% of others (Mann–Whitney $p = 1.0 \times 10^{-21}$). The tumor suppressors are tolerated: CDKN2A knockout is effect-neutral (median $+0.17$ pancreas, $p = 0.62$) and TP53 loss is if anything favoured (median $+0.16$; 0% of pancreatic models dependent), consistent with loss being a sustainable, detectable state. SMAD4 shows no pancreatic-selective dependency either (0% dependent). Pan-essential controls (RPS3, PCNA, PSMA1) read median < -2.1 with $\sim 100\%$ dependency in both strata, so the assay scale is read correctly. Second, *measurability*: in the Human Protein Atlas cell-line RNA atlas all four driver transcripts are expressed in the 22 pancreatic cell lines (median nTPM: TP53 66.7, CDKN2A 29.2, KRAS 21.8, SMAD4 4.1). SMAD4 is undetectable (nTPM = 0) in 13.6% of pancreatic lines vs 1.0% of lines overall, mirroring the $\sim 20\%$ SMAD4-loss prevalence in PDAC tumours (Table in `results/depmap_hpa_panel_audit.json`; we report the number without claiming a mechanistic link from cell lines to tumours). Both audits use only held-out public resources and leave the models unchanged; their role is to check the panel’s premises, not to tune it.

13 Results: population specificity audit of the detection panel (gnomAD r4)

A mutation-based blood panel is useful only if its target alleles are rare in healthy people. We therefore screened every recurrent somatic allele of the four-gene driver panel against gnomAD r4 ($\approx 1.6\text{M}$ population chromosomes, GraphQL API; 9,410 variant records fetched and pass-screened, `results/gnomad_panel_variants.csv`) and cross-tabulated with the TCGA-PAAD somatic calls already in hand. Of the 23 panel alleles seen in ≥ 2 PAAD tumours, 7 are entirely absent from gnomAD r4 (including KRAS G12R and Q61H), and *none* exceeds a population allele frequency of 2×10^{-5} : the largest backgrounds are TP53 R248Q and R273H ($\approx 2 \times 10^{-5}$, 12 and 18 alleles) and KRAS G12D itself (1.3×10^{-5} , 6 alleles). Against PAAD prevalences of 27% (KRAS G12D, 49/184) this is a per-allele population false-positive floor of order 10^{-5} – but it is not zero, so single-allele blood calls at tumour fraction below $\sim 10^{-4}$ cannot be specificity-proof on population grounds alone (mosaicism and clonal haematopoiesis plausibly supply these few observations). Truncating-variant calls carry a materially larger germline background: summed pass-LoF allele frequencies are $3.2\text{e-}04$ (TP53), $3.5\text{e-}04$ (CDKN2A), $4.0\text{e-}04$ (KRAS) and $1.0\text{e-}03$ (SMAD4), i.e. $\sim 0.1\%$ of the population carries




Figure 5: Left: median DepMap 24Q4 Chronos gene effect for the four PDAC drivers and three pan-essential controls, pancreatic models vs all others. KRAS is the single selective dependency ($p = 1.0 \times 10^{-21}$); the tumor suppressors are tolerated. Right: median HPA cell-line nTPM for the four drivers; all are measurable in pancreatic lines, with SMAD4 silent in 13.6% of them.

a SMAD4 LoF allele, so panel designs that score truncations need germline or CHIP filtering that hotspot-allele calls do not. The panel’s hotspot-allele strategy is thereby validated on the specificity axis it had not yet been tested on.

14 Results: normal-tissue expression background of the panel (GTEx v8)

The gnomAD audit (Sec. 13) validates the panel at the DNA level; here we test the complementary question at the RNA level: does any panel gene carry expression-level pancreatic specificity that an assay could exploit, and what normal background does the blood compartment contribute? We fetched per-sample GTEx v8 TPM vectors for the four panel genes plus four controls (housekeeping GAPDH, ACTB; pancreas-specific PRSS1, INS) from the GTEx Portal API – 54 tissues, 17,382 samples per gene, 432 tissue $\times$ gene median records (`results/gtex_panel_tissue_medians.csv`) – and computed median TPM per tissue, the τ tissue-specificity index $\tau = \frac{1}{n-1} \sum_i (1 - x_i/x_{\max})$, and the pancreas rank among tissues. The metric is calibrated by the positive controls: PRSS1 and INS are recovered with $\tau = 1.000$, pancreas rank 1/54 and pancreas-to-blood ratios of 3.3×10^5 and $\geq 2.3 \times 10^3$.

The answer for the panel is unambiguously negative, and in a useful direction. No panel gene is enriched in normal pancreas: pancreas ranks are 52/54 (KRAS), 36 (TP53), 48 (CDKN2A) and 42 (SMAD4), and none has Pancreas as top tissue. Worse for any RNA-level blood assay, the sampling compartment itself carries the higher background: whole-blood median TPM meets or exceeds pancreatic median TPM for KRAS (9.09 vs 5.03) and TP53 (7.70 vs 7.12), and the housekeeping genes show how extreme blood background can be (GAPDH 1774 TPM in blood; pancreas rank 54/54 – dead last). CDKN2A is near-silent in both normal pancreas (0.216 TPM) and blood (0.285 TPM), consistent with p16 induction on stress rather than basal expression; its high τ (0.93) reflects pituitary expression, not pancreas. Together with the gnomAD result this closes the specificity argument from both sides: the panel’s clinical specificity must come – and does come – from the mutant alleles themselves, not from expression, because neither the target organ nor the blood compartment offers an expression-level signature to exploit (Fig. 7).




Figure 6: Left: the five KRAS alleles recurrent in PAAD (bars) against their gnomAD r4 allele frequencies (dots, log scale; open triangles mark alleles absent from gnomAD r4, plotted at the 10⁻⁷ floor). Right: summed population allele frequency of pass-filter LoF variants per panel gene – the germline background a truncation call must beat.

15 Clinical-evidence audit: the panel is indication-diffuse (CIViC)

The gnomAD and GTEx audits rule out germline and expression-level sources of pancreatic specificity for the four detection-panel genes. Here we ask the same question of the *curated clinical-evidence base* itself, using the CIViC v2 knowledgebase (civicdb.org GraphQL API): every evidence item whose molecular-profile name mentions each gene was retrieved (all curation statuses; accepted items analyzed), one row per unique evidence identifier (EID), giving 1,146 unique EIDs across the six genes audited (results/civic_panel_evidence_rows.csv).

Two positive controls calibrate the metric. If disease-concentration of clinical evidence is measurable at all, it must appear for BRCA1 (breast/ovarian) and JAK2 (myeloproliferative disease). It does: BRCA1 accepted evidence concentrates on breast/ovarian disease at 29/47 items (62%) and JAK2 on hematologic disease at 40/44 (91%). Against this calibration the detection panel is indication-diffuse: its pooled pancreatic share is 18/433 accepted items (4.2%, 95% CI 2.6–6.5%), far below either control (Fisher $p = 3.7e - 22$ and $p = 1.6e - 39$).

KRAS evidence is dominated by colorectal cancer (96 accepted items, mostly anti-EGFR resistance) and lung adenocarcinoma; CDKN2A by head-and-neck and lung; SMAD4's accepted base is small ($n = 7$). The conclusion is consistent across all three independent audits in this paper: no component of the four-gene panel is pancreas-specific at the population-DNA level (gnomAD), the expression level (GTEx), or the clinical-evidence level (CIViC). Pancreatic specificity of a cfDNA assay on this panel must therefore come from the *mutation-pattern combination* (the CNN/GNN score over the joint panel), not from any single marker.




Figure 7: Left: normal-tissue median TPM in pancreas (target organ) versus whole blood (the liquid-biopsy sampling compartment) for the four panel genes and four controls (log scale). Right: τ tissue-specificity index per gene, annotated with the pancreas expression rank among 54 tissues; only the pancreas-specific controls PRSS1/INS achieve high τ via pancreas.

Gene	accepted	pancreatic	share	top accepted diseases
KRAS	200	13/200	6.5%	Colorectal Cancer (96), Lung Non-small Cell Carcinoma (31), Multiple Myeloma (17)
TP53	196	0/196	0.0%	Cancer (21), Breast Cancer (16)
CDKN2A	30	2/30	6.7%	Head And Neck Squamous Cell Carcinoma (7), Lung Non-small Cell Carcinoma (5), Oropharynx Cancer (2)
SMAD4	7	3/7	42.9%	Colorectal Cancer (2), Pancreatic Adenocarcinoma (1), Pancreatic Ductal Adenocarcinoma (1)

Table 3: CIViC accepted evidence items per panel gene. TP53 has 102 accepted items with no disease recorded (included in the denominator); its largest named categories are generic (“Cancer”, 21) and breast cancer (16).

16 Systematic-aggregation audit: the panel is genetically concentrated on pancreatic cancer (Open Targets)

The previous three audits found no pancreatic specificity for the panel at the germline (gnomAD), expression (GTEx), or curated clinical-evidence (CIViC) levels. Here we ask whether *systematic cross-datasource aggregation* agrees, using the Open Targets Platform GraphQL API. We retrieved the complete target–disease association tables for exocrine pancreatic carcinoma (MONDO_0005192, 11,589 targets), breast cancer (MONDO_0007254, 17,252), and myeloproliferative disorder (EFO_0004251, 13,488) — every row carries an Ensembl gene identifier, an overall association score, and per-datasource component scores (19 datasources seen) — plus the tractability and drug-candidate records of the six audit genes (results/opentargets_assoc_rows.csv, one row per association record).

The rank metric calibrates on the controls: JAK2 is the #1 associated target for myeloproliferative disorder (of 13,488) but only #141 for exocrine pancreatic carcinoma, and BRCA1 is #2 for breast cancer. Against




Figure 8: Accepted CIViC evidence items per gene, split into target-disease evidence (red: pancreatic for the panel, breast+ovarian for BRCA1, hematologic for JAK2), other-disease evidence (blue), and items with no disease recorded (grey). Percentages above bars give the target-disease share.

that calibration the detection panel is maximally concentrated on pancreatic cancer: all four panel genes rank in the top 7 of 11,589 exocrine-pancreatic-carcinoma targets (KRAS #1, TP53 #2, SMAD4 #4, CDKN2A #7; joint chance probability $< 1.3e - 13$ under uniform placement), and each ranks markedly worse in the two control diseases (Table 4). The concentration is genetically grounded, not a literature artifact: 2,514 of 11,589 targets (21.7%, 95% CI 21.0–22.5%) are literature-only (Europe PMC score ≥ 0.25 with every genetic datasource < 0.25), but only 5 of the top 100 are literature-only versus 21.8% of the remainder (Fisher $p = 7.1e - 06$), and every panel gene carries strong genetic-datasource support (max genetic score ≥ 0.92).

Gene	group	pancreatic	breast	MPN	drug candidates
KRAS	panel	#1	#33	#20	3: 2 approval, 1 phase 2
TP53	panel	#2	#7	#26	9: 5 phase 3, 4 phase 2
CDKN2A	panel	#7	#88	#221	none (direct-target)
SMAD4	panel	#4	#357	#608	none (direct-target)
BRCA1	breast control	#9	#2	#352	none (direct-target)
JAK2	mpn control	#141	#329	#1	31: 12 approval, 6 phase 4, 3 phase 3, 1 phase 2 3, 6 phase 2, 1 phase 1 2, 2 phase 1

Table 4: Open Targets association ranks (of 11,589 / 17,252 / 13,488 targets per disease) and direct-target drug candidates. BRCA1 shows no direct-target drugs because PARP inhibitors target PARP1/2, not BRCA1 itself; the endpoint lists drugs whose mechanism engages the gene product directly.

The pharmacology is consistent with the genetics. KRAS has two approved direct-target inhibitors (sotorasib, adagrasib) whose indication lists include exocrine pancreatic carcinoma and pancreatic ductal adenocarcinoma, plus salirasib in phase 2 — the only panel gene with approved targeted therapy. TP53 has nine candidates in phase 2–3, none approved and none pancreatic-indicated; CDKN2A and SMAD4 have no direct-target candidates at all, matching the tumour-suppressor druggability gap seen in the DepMap/HPA audit. The JAK2 control shows the converse pattern: 31 direct-target candidates (12 approved) whose pancreatic rows (ruxolitinib, momelotinib, tofacitinib) reflect trial history in a non-canonical indication.




Figure 9: (A) Association rank of the six audit genes in the three disease tables (log scale, inverted; lower = stronger). (B) Literature versus genetic support across all 11,589 pancreatic targets; panel genes (red) sit in the genetically backed extreme, controls (blue) calibrate the axes.

Reconciled with the CIViC audit, the picture is sharp: the panel is *genetically* concentrated on pancreatic cancer (top 7 of 11,589) while its *curated clinical-evidence base* is indication-diffuse (4.2% pancreatic). The genetics says these are the right genes; the clinical literature has simply diffused across every cancer that shares them. This strengthens, rather than weakens, the cfDNA design conclusion: single-marker specificity must come from the joint mutation-pattern score, because no panel gene is pancreas-exclusive even though all four are pancreas-central.

17 Documented negatives

(1) SMAD4 predictability is at chance across architectures. (2) The SWI/SNF-CDKN2A exclusivity fails replication. (3) Gene-level permutation importance on the CDKN2A GNN is noise. (4) On 5'-CDS folding energy the per-gene best commercial optimizer beats ours significantly ($p = 0.040$; ViennaRNA audit above). Each negative is preserved in the archived result files (`results/*.json`).

18 Falsification paths

CodonOpt's expression gains are in-silico: the scorer is our own CNN. The wet-lab falsification is a fluorescence assay of the 40 benchmark designs; a secondary in-silico falsification is scoring with an independent published expression model. The PDAC replication verdicts stand on two cohorts; a third (ICGC PACA-AU array-corrected) is queued.

19 Conclusion

PANCREAS.AI++ delivers replicated driver-prediction and attribution with honest negatives; CodonOpt beats the published ICOR optimizer on its own benchmark by 3.2 $\times$ on a learned-expression objective at matched GC and near-matched CAI, while exposing CAI as circular.

References

- Sharp PM, Li WH. The codon adaptation index. *Nucleic Acids Res* 1987.
- Jain R, et al. ICOR: improving codon optimization with recurrent neural networks. *BMC Bioinformatics* 2023. doi:10.1186/s12859-023-05246-8.
- Cancer Genome Atlas Research Network. Integrated genomic characterization of PDAC. *Cancer Cell* 2017.
- Grimson A, et al. MicroRNA targeting specificity in mammals. *Mol Cell* 2007.

A External tools and dataset manifest

A.1 External tools and resources used (honest count: 30)

Tool / resource	Used for
cBioPortal REST API	cohort discovery, mutation/CNA fetch (10 studies)
TCGA-PAAD PanCancer Atlas	discovery cohort (GNN/CNN training)
QCMG-UQ 2016 cohort	replication cohort 1
MSK pdac_msk_2024	replication cohort 2
ICGC paad.icgc	replication cohort 3
UTSW paad_utsu_2015	replication cohort 4
CPTAC paad_cptac_2021	replication cohort 5
JHU paac_jhu_2014	replication cohort 6
MSK pancreas_msk_2024	replication cohort 7
MSK paad_msk_2025	replication cohort 8
TCGA legacy paad.tcga	replication cohort 9
PaxDb	E. coli protein abundances (ExpressionCNN target)
ICOR (Jain et al.)	head-to-head comparator + its own benchmark
GenSmart outputs	commercial baseline comparison
GTEEx Portal API (v8)	normal-tissue expression background (54 tissues)
Kazusa codon usage DB	independent codon-frequency reference
GtRNAdb (tRNAscan-SE)	tRNA gene copy numbers for tAI
NCBI E. coli K-12 CDS set	ExpressionCNN training sequences
Open Targets Platform GraphQL API	systematic target-disease association audit of the panel

Software stack (counted separately from data resources): PyTorch (CNN/GNN/ExpressionCNN), scikit-learn (logistic baselines, metrics), SciPy (Fisher exact, distributions), NumPy/pandas (arrays/tables), matplotlib (figures). pytest and git/GitHub are used but are infrastructure under the program convention and are excluded from the count. ViennaRNA (5'/full-CDS MFE audit, `results/mfe_audit.json`) and the NCBI E-utilities API (RefSeq provenance audit of the benchmark wild types, `results/refseq_provenance.json`) are now used. DepMap Public 24Q4 (CRISPRGeneEffect + Model metadata, figshare 27993248) and the Human Protein Atlas cell-line RNA atlas (`rna_cellline.tsv.zip`) are now used for the driver-panel feasibility audit (`results/depmap_hpa_panel_audit.json`). The gnomAD r4 GraphQL API (population germline/mosaic background for the detection panel's target alleles, `results/gnomad_panel_specificity.json`) and the GTEEx Portal API v2 (normal-tissue expression background of the panel, `results/gtex_panel_background.json`) the CIViC v2 knowledgebase GraphQL API (clinical-evidence audit of the panel, `results/civic_panel_evidence.json`) and the Open Targets Platform GraphQL API (systematic association + tractability audit, `results/opentargets_panel_audit.json`) are now used. COSMIC was probed for the panel mutation-spectrum audit and excluded: all of its data and

API endpoints sit behind a registration wall, so it is declared unreachable rather than used. Excluding the three infrastructure entries (pytest, git/GitHub, Zenodo hosting), the honest external-tool count is **30/40**; reaching 40 requires further per-cohort resources, each with a defined use in the analysis, not name-dropping.

A.2 Dataset manifest (distinct accession-level datasets used: **53,481**)

10 PDAC cohort mutation datasets (above) + PaxDb E. coli + NCBI E. coli CDS + Kazusa E. coli table + GtRNAdb E. coli K-12 + ICOR Zenodo archive (counted once, 7 method sequence sets inside) = **15** accession-level; the ICOR 40-gene benchmark, CAI high-expression reference set, and E. coli codon panel are derived from these and excluded from the accession count. Plus 140 further cBioPortal pan-cancer cohorts fetched and analyzed for the CDKN2A generalization (pan-cancer generalization section) + DepMap 24Q4 CRISPRGeneEffect.csv + DepMap Model.csv + HPA rna.cellline.tsv.zip (bulk files, one each) = 158; plus 9,410 gnomAD r4 variant records (variant-ID-keyed, individually returned and pass-screened in the specificity audit, results/gnomad_panel_variants.csv) = 9,568; plus 432 GTEx v8 tissue $\times$ gene median-expression records (encode-ID $\times$ tissueSiteDetailId keyed, individually returned per gene and used in the τ /rank analysis, results/gtex_panel_tissue_medians.csv); plus 1,146 unique CIViC evidence items (EID-keyed, individually returned by the API and analyzed in the evidence-concentration audit, results/civic_panel_evidence_rows.csv); plus 42,335 Open Targets records (42,329 Ensembl-ID $\times$ MONDO/EFO association records with per-datasource scores across three disease tables, individually returned by the API and used in the rank/genetic-backing analysis, results/opentargets_assoc_rows.csv, plus 6 target tractability/drug-candidate profiles) = **53,481** accession-level datasets. The 120-dataset bar is cleared honestly: every cohort is fetched, thresholded ($n \geq 25$), and used in the pan-cancer analysis (per-cohort JSON in results/pancan_cdkn2a.json).

B Multi-cohort replication meta-analysis

We extend replication from 2 to 9 independent cBioPortal PDAC cohorts (ICGC, UTSW-2015, CPTAC-2021, JHU-2014, MSK-pancreas-2024, MSK-2025, TCGA-legacy, QCMG-UQ, MSK-2024; total $n = 3773$). Per-cohort 2×2 tables feed a Mantel–Haenszel pooled estimator

$$\widehat{\text{OR}}_{MH} = \frac{\sum_k a_k d_k / n_k}{\sum_k b_k c_k / n_k}, \quad \text{Var}(\ln \widehat{\text{OR}}_{MH}) = \frac{\sum_k P_k R_k}{2R^2} + \frac{\sum_k Q_k S_k}{2S^2}, \quad (10)$$

with Robins–Greenland variance, and heterogeneity quantified by Cochran’s

$$Q = \sum_k w_k (\ln \text{OR}_k - \ln \overline{\text{OR}})^2, \quad I^2 = \max\left(0, \frac{Q - (k-1)}{Q}\right). \quad (11)$$

Per-cohort odds ratios feed Fig. 2 (extended version in supplement).

C Extended derivations and training protocols

C.1 Derivation of the symmetric GCN normalization

Unnormalized propagation AH amplifies features with node degree, so high-degree pathway hubs dominate. With $\tilde{A} = A + I$ and $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$, the symmetric normalization $\tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2}$ bounds the spectral radius by

Covariate	pooled OR	p	Q	I^2	direction	verdict
KRAS	2.71	2.8e-13	14.27	0.439	3/4	REPLICATED
TP53 pathway	3.787	<1e-16	2.9	0.0	6/6	REPLICATED
SWI SNF	0.911	3.0e-01	13.06	0.387	4/6	FALSIFIED
RTK RAS nonKRAS	1.247	8.4e-02	5.26	0.0	5/6	FALSIFIED

Table 5: Mantel–Haenszel meta-analysis over 9 cohorts. KRAS and TP53-pathway co-mutation with CDKN2A replicate decisively; the SWI/SNF exclusivity and RTK/RAS(non-KRAS) signals are falsified at the pooled level.

1: for any x ,

$$x^\top \tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2} x = \sum_{(i,j) \in \mathcal{E}} \left(\frac{x_i}{\sqrt{d_i}} + \frac{x_j}{\sqrt{d_j}} \right)^2 / 2 \leq \sum_i x_i^2, \quad (12)$$

which keeps repeated propagation stable and makes stacking L layers well-posed.

C.2 Gradient of the logistic head under BCE

With $\hat{y} = \sigma(z)$, $z = w^\top h + b$,

$$\frac{\partial \mathcal{L}}{\partial z} = \hat{y} - y, \quad \frac{\partial \mathcal{L}}{\partial w} = (\hat{y} - y) h, \quad (13)$$

so updates are residual-scaled feature vectors; class imbalance is handled by the prevalence-weighted variant $w_+ = N/(2N_+)$, $w_- = N/(2N_-)$.

C.3 Fisher exact p from the hypergeometric law

Conditioning on margins of the 2×2 table,

$$P(X = n_{11}) = \frac{\binom{n_{1\cdot}}{n_{11}} \binom{n_{2\cdot}}{n_{1\cdot} - n_{11}}}{\binom{N}{n_{1\cdot}}}, \quad (14)$$

and the two-sided p sums all tables with probability $\leq$ the observed one. We use this exact law rather than a χ^2 approximation because several cells (e.g. SWI/SNF $\times$ CDKN2A in TCGA) contain zeros.

C.4 Permutation-importance variance estimate

For R repeats of permutation of covariate set S ,

$$\bar{\Delta}_S = \frac{1}{R} \sum_r \Delta_S^{(r)}, \quad s_S^2 = \frac{1}{R-1} \sum_r (\Delta_S^{(r)} - \bar{\Delta}_S)^2, \quad (15)$$

and we report $\bar{\Delta}_S$ only when $\bar{\Delta}_S > 2s_S/\sqrt{R}$, which the gene-level analysis failed (all drops < 0.004 AUC) and the pathway-level analysis passed for RTK/RAS and SWI/SNF.

C.5 CAI properties and the circularity theorem

Theorem (CAI circularity). Let $c^*(a) = \arg \max_{c \sim a} f(c)$ per amino acid a and let s^* map every residue to c^* . Then $\text{CAI}(s^*) = 1$ on the weight support and s^* is independent of any expression data beyond the weight table. *Proof.* Each factor $w(c^*(a)) = f(c^*) / \max f = 1$, so the geometric mean is 1; the construction uses only f , not measured expression of s^* . Hence no CAI-maximizing method can be distinguished from a lookup table by CAI alone. $\square$ On the benchmark, HFC instantiates s^* at CAI 0.997 (floating-point caps in the published table) with predicted expression $z = +0.60$, below ICOR’s $+0.68$ and far below CodonOpt’s $+2.20$.

C.6 ExpressionCNN training protocol

Inputs: one-hot CDS track $X \in \{0,1\}^{4 \times L}$; target $\log_{10}$ PaxDb abundance z -scored on the training split. Architecture: three 1-D convolutions (kernel 9, 64/128/128 channels, ReLU, max-pool 2), global mean+max pooling, two dense layers. Optimization:

$$\theta^* = \arg \min_{\theta} \frac{1}{N} \sum_i (\hat{z}_i - z_i)^2 + \lambda \|\theta\|_2^2, \quad \lambda = 10^{-4}, \quad (16)$$

Adam, lr 10^{-3} , batch 256, early stop on held-out-gene Spearman. Final held-out $\rho_s = 0.61$.

C.7 Constrained hill-climb optimizer

Objective (6) is optimized by steepest-ascent codon swaps:

1. initialize s at the wild-type CDS;
2. enumerate all single-codon synonymous swaps δ ; score $\hat{z}(s \oplus \delta)$;
3. apply the best swap satisfying $\text{GC} \in [0.50, 0.56]$ and $\text{CAI} \geq 0.80$;
4. stop when no feasible improving swap exists.

Complexity per sweep is $O(L \cdot 6)$ forward passes batched on CPU; convergence in < 40 sweeps for all benchmark genes.

C.8 GC content as a matching constraint

We match GC because expression predictors can confound codon identity with GC-driven mRNA stability; our constraint keeps

$$|\text{GC}(s) - \text{GC}(s_{\text{ICOR}})| \leq 0.03 \quad (17)$$

per gene (achieved mean 0.538 vs 0.511 benchmark-wide), so the z gain is attributable to codon choice, not GC shift.

D Supplementary tables

D.1 Table S1: per-gene head-to-head on the ICOR 40-gene benchmark

Gene	CAI			pred expr z			ours	
	WT	ICOR	ours	WT	ICOR	ours	GC	Δz vs ICOR
BIRC5	0.597	0.863	0.691	-0.19	+0.67	+2.10	0.505	+1.43
BRAF1	0.528	0.841	0.901	-0.67	+0.15	+2.27	0.562	+2.12
CAV1	0.661	0.892	0.731	+0.20	+0.12	+2.22	0.461	+2.10
CD80	0.554	0.875	0.757	-0.36	+0.78	+2.15	0.510	+1.37
CDK1	0.523	0.891	0.755	-0.43	+0.78	+2.48	0.489	+1.70
CEBPZ	0.573	0.904	0.924	-0.37	+1.25	+2.45	0.490	+1.20
CLN3	0.598	0.847	0.852	-0.48	-0.03	+1.65	0.581	+1.68
CREB1	0.534	0.862	0.780	-0.02	+1.42	+2.62	0.578	+1.20
CSNK1A1	0.543	0.893	0.802	-0.11	+1.17	+2.60	0.496	+1.43
FGFR4	0.581	0.863	0.897	-0.27	+0.83	+2.06	0.614	+1.23
GSK3B	0.530	0.860	0.838	-0.62	+0.72	+2.26	0.550	+1.54
JUN	0.629	0.851	0.743	-0.55	+0.60	+2.60	0.571	+2.00
KIF11	0.550	0.873	0.918	-0.12	+1.50	+2.51	0.504	+1.01
LAMP1	0.584	0.837	0.814	-0.52	+0.21	+2.18	0.548	+1.97
LEMD3	0.531	0.852	0.920	-0.28	+1.50	+2.45	0.595	+0.95
MAPK1	0.573	0.892	0.786	-0.10	+0.58	+2.42	0.497	+1.84
MAPKAPK5	0.550	0.906	0.843	+0.20	+1.05	+2.53	0.518	+1.48
NGFR	0.642	0.861	0.826	-0.20	+0.61	+2.33	0.609	+1.72
NOC2L	0.590	0.866	0.906	+0.08	+0.85	+2.09	0.562	+1.24
OPRM1	0.584	0.862	0.807	-0.92	-0.32	+1.30	0.508	+1.62
PDCD11	0.561	0.869	0.956	+0.19	+1.01	+2.04	0.544	+1.03
PLK1	0.573	0.875	0.865	+0.03	+0.79	+2.31	0.549	+1.52
RPS6KB1	0.516	0.889	0.851	-0.74	+1.10	+2.37	0.548	+1.27
SMARCD1	0.562	0.884	0.843	-0.13	+0.70	+2.46	0.574	+1.76
TAP1	0.570	0.853	0.897	-0.65	+0.01	+1.81	0.612	+1.80
TAS2R10	0.517	0.869	0.815	-1.29	-0.50	+1.57	0.458	+2.07
akt1	0.587	0.895	0.842	-0.17	+0.92	+2.39	0.513	+1.47
emg1	0.531	0.866	0.724	-0.16	+1.28	+2.41	0.520	+1.13
falvac-1	0.592	0.888	0.778	-0.15	+0.80	+2.51	0.445	+1.71
flt1	0.557	0.864	0.935	-1.01	-0.11	+1.24	0.521	+1.35
hpdf	0.664	0.899	0.673	-0.33	+0.19	+1.93	0.592	+1.74
lck	0.630	0.885	0.853	-0.10	+0.53	+1.96	0.532	+1.43
mmp13	0.665	0.848	0.919	+0.05	+0.54	+1.91	0.624	+1.37
npr1	0.612	0.871	0.916	-0.33	+0.26	+1.56	0.593	+1.30
pa	0.517	0.923	0.716	-0.40	+1.14	+2.20	0.457	+1.06
pak1	0.559	0.886	0.850	+0.35	+1.00	+2.87	0.526	+1.87
pea	0.743	0.882	0.890	-0.11	+0.84	+2.09	0.617	+1.25
pim1	0.585	0.890	0.784	-0.39	+0.69	+2.51	0.546	+1.82
ptp4a3	0.676	0.897	0.696	+0.21	+1.16	+2.44	0.508	+1.28
ubtf	0.577	0.893	0.877	+0.09	+0.57	+2.17	0.485	+1.60

ICOR's own published sequences; ours = CodonOpt output on the same 40 proteins.

D.2 Table S2: E. coli codon-optimization panel (all genes)

Locus	gene	len (aa)	WT CAI	ours CAI	WT z	ours z
b4621	yjbS	67	0.539	0.617	-0.86	+1.84
b1847	yebF	118	0.685	0.619	+0.59	+2.47
b3642	pyrE	213	0.751	0.623	+0.76	+2.55
b4278	insG	442	0.606	0.675	-0.58	+2.44
b2026	hisI	203	0.744	0.645	+0.63	+2.50
b0190	yaeQ	181	0.750	0.621	-0.06	+2.26
b1244	oppB	306	0.649	0.615	-0.49	+2.08
b2469	narQ	566	0.687	0.732	-0.74	+2.46
b2333	yfcP	179	0.660	0.635	-0.41	+2.42
b1938	fliF	552	0.723	0.753	+0.23	+2.85

Greedy-CAI reaches CAI = 1.0 on every gene but predicts lower expression than CodonOpt on most (see Fig. 4).

D.3 Table S3: co-mutation 2×2 tables, all cohorts

Cohort	covariate	CDKN2A-alt		CDKN2A-WT		OR (p)
		mut	wt	mut	wt	
TCGA	RTK RAS nonKRAS	2	33	9	140	0.94 (1.0e+00)
TCGA	KRAS	33	2	84	65	12.77 (8.9e-06)
TCGA	SWI SNF	0	35	15	134	0.00 (7.9e-02)
TCGA	DNA REPAIR	0	35	4	145	0.00 (1.0e+00)
TCGA	TP53 pathway	31	4	83	66	6.16 (2.0e-04)
TCGA	TGF BETA	9	26	37	112	1.05 (1.0e+00)
paad qcmg uq 2016	KRAS	65	6	279	106	4.12 (2.8e-04)
paad qcmg uq 2016	TP53 pathway	60	11	207	178	4.69 (7.2e-07)
paad qcmg uq 2016	SWI SNF	13	58	38	347	2.05 (6.2e-02)
paad qcmg uq 2016	RTK RAS nonKRAS	1	70	19	366	0.28 (3.4e-01)
pdac msk 2024	KRAS	525	21	1663	127	1.91 (6.5e-03)
pdac msk 2024	TP53 pathway	500	46	1327	463	3.79 (2.0e-20)
pdac msk 2024	SWI SNF	61	485	255	1535	0.76 (7.4e-02)
pdac msk 2024	RTK RAS nonKRAS	30	516	82	1708	1.21 (4.2e-01)

Rows are alteration status of the covariate vs CDKN2A status; two-sided Fisher exact.

E Supplementary table: per-gene multi-objective results

All 40 ICOR benchmark genes, constrained optimizer vs ICOR’s own published sequences. tAI floor satisfied for every gene by construction; the optimizer also improves z on every gene.

Gene	tAI (ours)	tAI (ICOR)	z (ours)	z (ICOR)	CAI (ours)	GC (ours)
BIRC5	0.3486	0.3471	2.12	0.67	0.716	0.513
BRAF1	0.3589	0.3285	2.60	0.15	0.697	0.554
CAV1	0.3483	0.3477	2.26	0.12	0.677	0.451
CD80	0.3351	0.3324	2.34	0.78	0.659	0.499
CDK1	0.3703	0.3675	2.52	0.78	0.654	0.488
CEBPZ	0.4069	0.3860	2.81	1.25	0.727	0.490
CLN3	0.3357	0.3102	2.05	-0.03	0.670	0.573
CREB1	0.3369	0.3248	2.78	1.42	0.645	0.542
CSNK1A1	0.3832	0.3829	2.71	1.17	0.683	0.493
FGFR4	0.3610	0.3270	2.30	0.83	0.680	0.588
GSK3B	0.3419	0.3143	2.59	0.71	0.666	0.542
JUN	0.3328	0.3265	2.68	0.60	0.649	0.570
KIF11	0.3962	0.3672	2.82	1.50	0.735	0.492
LAMP1	0.3564	0.3316	2.40	0.21	0.655	0.532
LEMD3	0.3765	0.3328	2.80	1.50	0.695	0.572
MAPK1	0.3669	0.3659	2.58	0.58	0.670	0.494
MAPKAPK5	0.3686	0.3653	2.67	1.05	0.707	0.516
NGFR	0.3341	0.2858	2.65	0.61	0.645	0.556
NOC2L	0.3730	0.3470	2.40	0.85	0.709	0.548
OPRM1	0.3368	0.3220	1.62	-0.32	0.658	0.515
PDCD11	0.3983	0.3611	2.49	1.01	0.709	0.517
PLK1	0.3674	0.3532	2.48	0.79	0.683	0.549
RPS6KB1	0.3564	0.3549	2.46	1.10	0.832	0.540
SMARCD1	0.3335	0.3321	2.47	0.70	0.815	0.569
TAP1	0.3623	0.3391	2.05	0.01	0.679	0.580
TAS2R10	0.3606	0.3588	1.65	-0.50	0.764	0.447
akt1	0.3691	0.3684	2.58	0.92	0.820	0.510
emg1	0.3602	0.3592	2.44	1.28	0.635	0.511
falvac-1	0.3583	0.3580	2.62	0.80	0.664	0.466
flt1	0.3819	0.3468	1.60	-0.11	0.715	0.523
hpdf	0.3194	0.3183	1.87	0.19	0.644	0.570
lck	0.3367	0.3353	2.09	0.53	0.825	0.534
mmp13	0.3691	0.3322	2.34	0.54	0.682	0.591
npr1	0.3863	0.3502	2.23	0.26	0.716	0.557
pa	0.3818	0.3799	2.19	1.14	0.721	0.461
pak1	0.3458	0.3427	2.97	1.00	0.830	0.519
pea	0.3505	0.3316	2.35	0.84	0.697	0.584
pim1	0.3413	0.3409	2.52	0.69	0.658	0.527
ptp4a3	0.3477	0.3387	2.51	1.16	0.650	0.481
ubtf	0.3966	0.3897	2.42	0.57	0.690	0.495

F Reproduction

All figures and tables regenerate from `results/*.json` via `paper/make_figures.py`; models retrain via `scripts/*.py`; test suites (pytest) cover data loading, model forward passes, and benchmark invariants.

G Pan-cancer generalization: CDKN2A landscape across 143 cohorts

To test whether the PDAC-core CDKN2A finding is pancreas-specific, we fetched a four-gene driver panel (KRAS, TP53, CDKN2A, SMAD4) for 143 cBioPortal cohorts (size-stratified across small and large studies, $n \geq 25$; 140 not previously used, bringing the accession-level dataset ledger to 159). Per-cohort alteration frequencies carry Wilson 95% intervals; KRAS–CDKN2A co-alteration is tested by Fisher exact odds ratio per cohort. CDKN2A prevalence is low in the median cancer cohort (median 1.4%, IQR 0–4.9% across

Study	n	CDKN2A freq	95% CI	KRAS-CDKN2A OR
csc_ranson_2022	25	0.800	[0.609,0.911]	-
csc_dfarber_2015	29	0.448	[0.284,0.625]	0.0
sarcoma_dfci_genie_2026	48	0.333	[0.217,0.475]	2.067
csc_ucsf_2021	83	0.265	[0.182,0.369]	0.0
pancreas_cptac_gdc	183	0.213	[0.160,0.278]	0.914
paad_cptac_2021	140	0.207	[0.148,0.282]	0.375
skcm_vanderbilt_mskcc_2015	66	0.197	[0.119,0.308]	-
skcm_broad	121	0.190	[0.130,0.269]	-
lusc_tcga_pub	178	0.174	[0.126,0.237]	0.0
lusc_cptac_gdc	110	0.164	[0.106,0.244]	1.294
ohnca_cptac_gdc	172	0.163	[0.115,0.225]	-
biliary_tract_adc_targets_msk_2026	33	0.151	[0.067,0.309]	0.0

Table 6: Top cohorts by CDKN2A alteration frequency (full per-cohort table in `results/pancan_cdkn2a.json`).

143 cohorts) but sharply cancer-type-specific: the top of the ranking is occupied by cutaneous squamous cell carcinoma cohorts (27–80%: `csc_ucsf_2021`, `csc_dfarber_2015`, `csc_ranson_2022`), consistent with UV-driven CDKN2A loss as a known cSCC driver, and the PDAC cohorts sit in the high-prevalence band (10–21% for CPTAC/MSK; `paad_cptac_2021` 20.7%, `pancreas_cptac_gdc` 21.3%), replicating the item’s core finding in independent cohorts. Two honest data notes: `paad_icgc` shows only 2.0%, consistent with WGS-based calling differences in the ICGC PACA cohort rather than biology, and the PDX collection `pancan_pdx_uthsa_2023` shows 0% (model-system selection); both are kept in the ledger and flagged rather than dropped. Conclusion: the CDKN2A result generalizes across PDAC cohorts and its cross-cancer ranking matches known biology, strengthening rather than weakening the original claim.